

HADOOP TECHNOLOGY

What is Hadoop Technology??

- The most well known technology used for Big Data is Hadoop.
- It is actually a large scale batch data processing system

Why Hadoop ??

- Distributed cluster system
- Platform for massively scalable applications
- Enables parallel data processing

Developers of Hadoop Technology:



Michael j. cafarella



Doug cutting

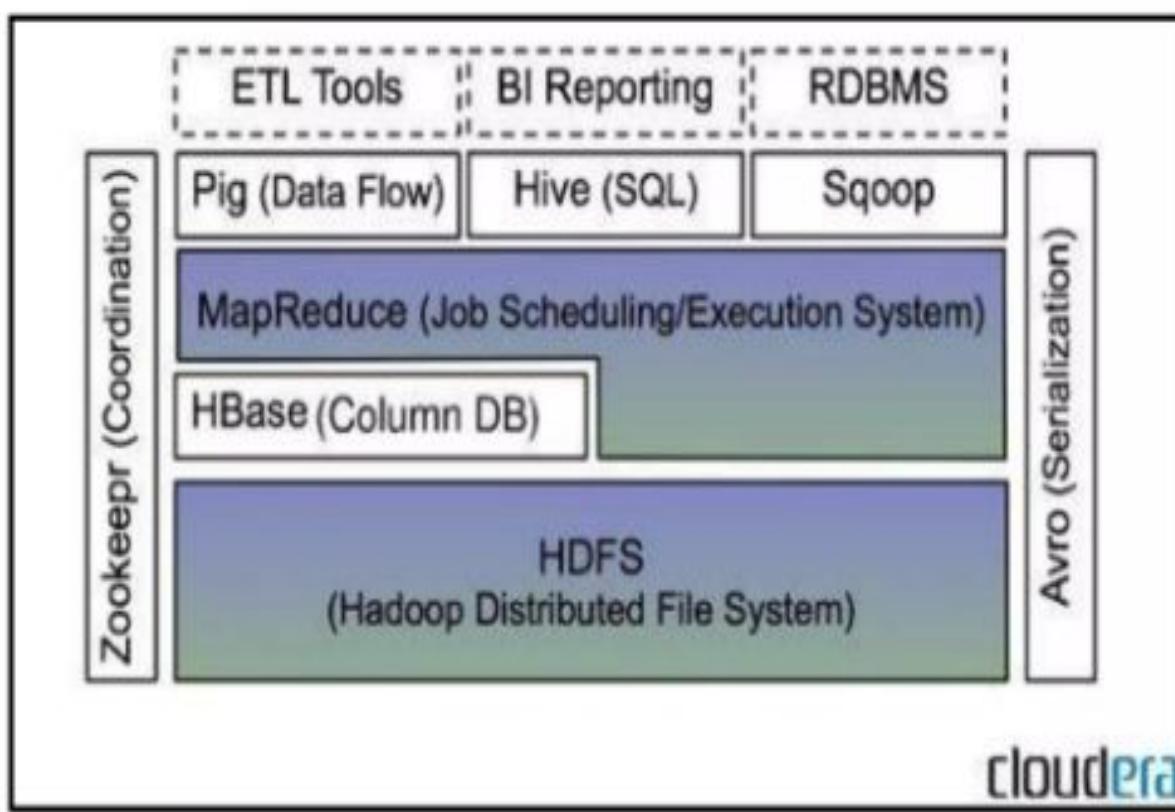
Famous Hadoop users



Hadoop Features

- Hadoop provides access to the file systems
 - The Hadoop Common package contains the necessary JAR files and scripts
 - The package also provides source code,
documentation and a contribution section that includes projects from the Hadoop Community.

HADOOP ARCHITECTURE



cloudera

Core-Components of Hadoop:



Hadoop distributive file system.

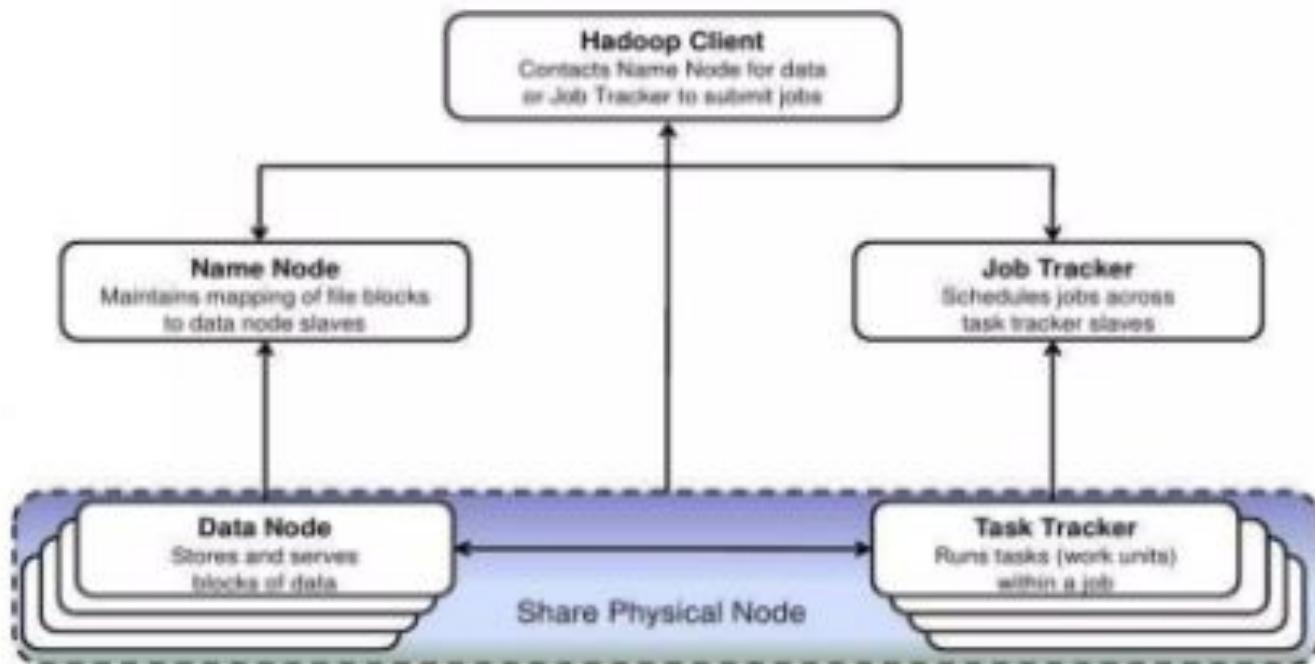


Map reduce.

What is HDFS ?

- Distributed file system
- Traditional hierarchical file organization
- Single namespace for the entire cluster
- Write-once-read-many access model
- Aware of the network topology

Hadoop High Level Architecture



Hadoop cluster

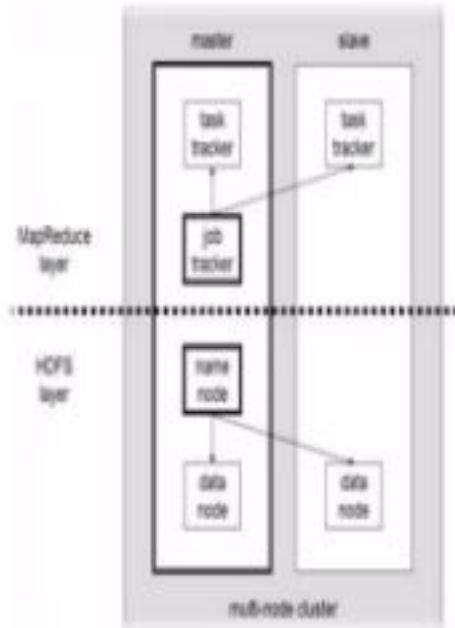
- A Small Hadoop Cluster Include a single master & multiple worker nodes

Master node:

Data Node
Job Tracker
Task Tracker
Name Node

Slave node:

Data Node
Task Tracke

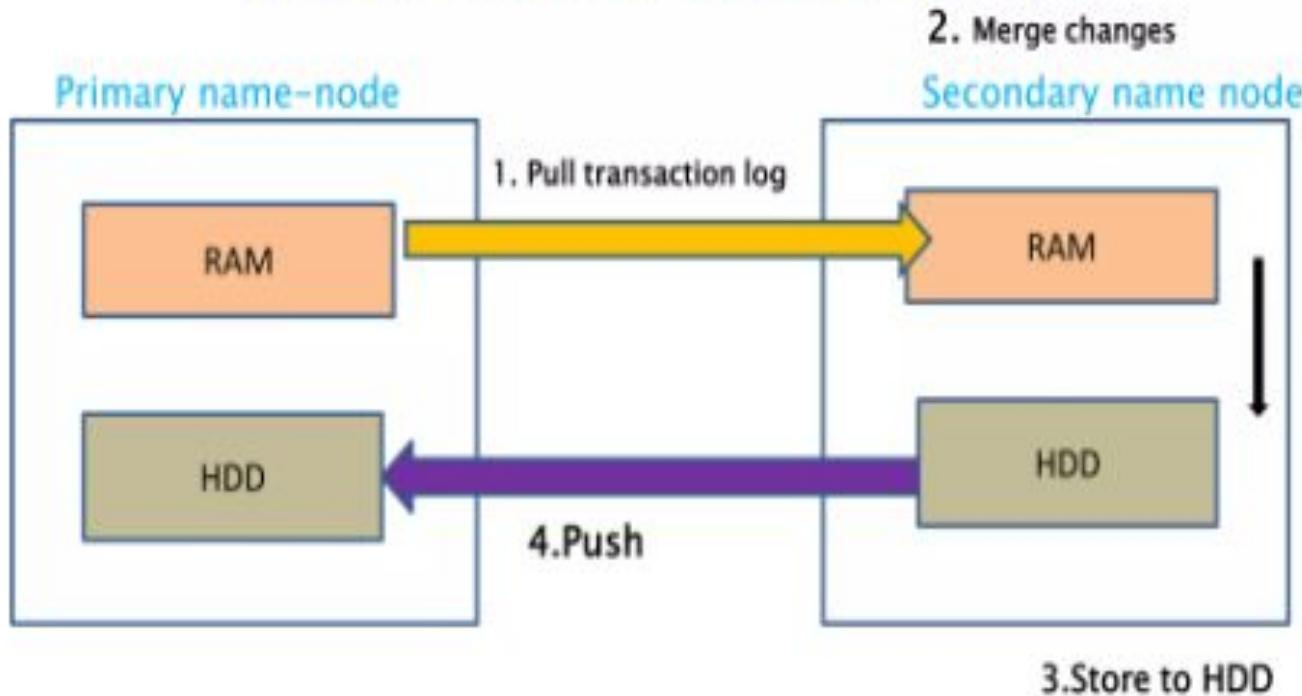


HDFS – Name Node Features

Metadata in main memory:

- List of files
- List of blocks for each file
- List of Data Nodes for each block
- File attributes
- Creation time
- Records every change in the metadata

HDFS-name node architecture



HDFS-Data node

- Block Server Stores data in the local file system
- Periodic validation of checksums
- Periodically sends a report of all existing blocks to the Name Node



Hadoop MAPREDUCE

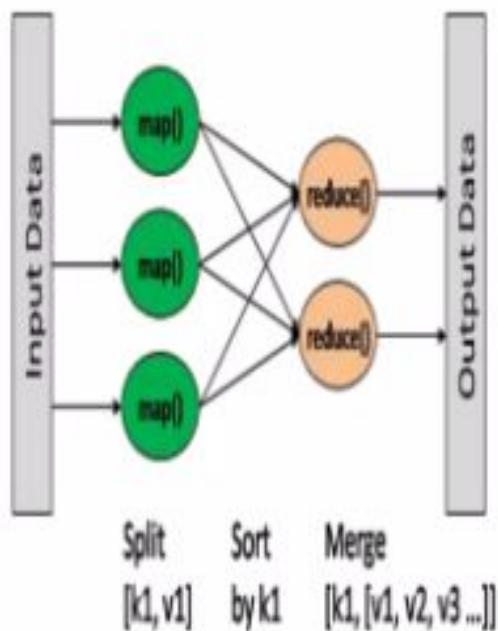
Map reduce implementation:

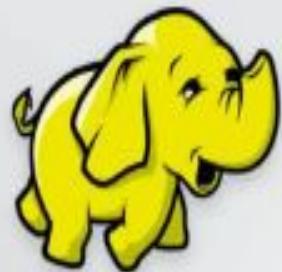
Job Tracker:

Splitting into map and reduce tasks
Scheduling tasks on a cluster node

Task Tracker:

Runs Map Reduce tasks periodically





What is Hadoop?

- Hadoop is an open source software programming framework for storing a large amount of data and performing the computation.
- Its framework is based on Java programming with some native code in C and shell scripts.
- Hadoop is used for some advanced level of analytics, which includes Machine Learning and data mining

Need for Hadoop

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Programmers

No longer need to worry about



Q: Where file is located?

Q: How to handle failures & data lost?

Q: How to divide computation?

Q: How to program for scaling?

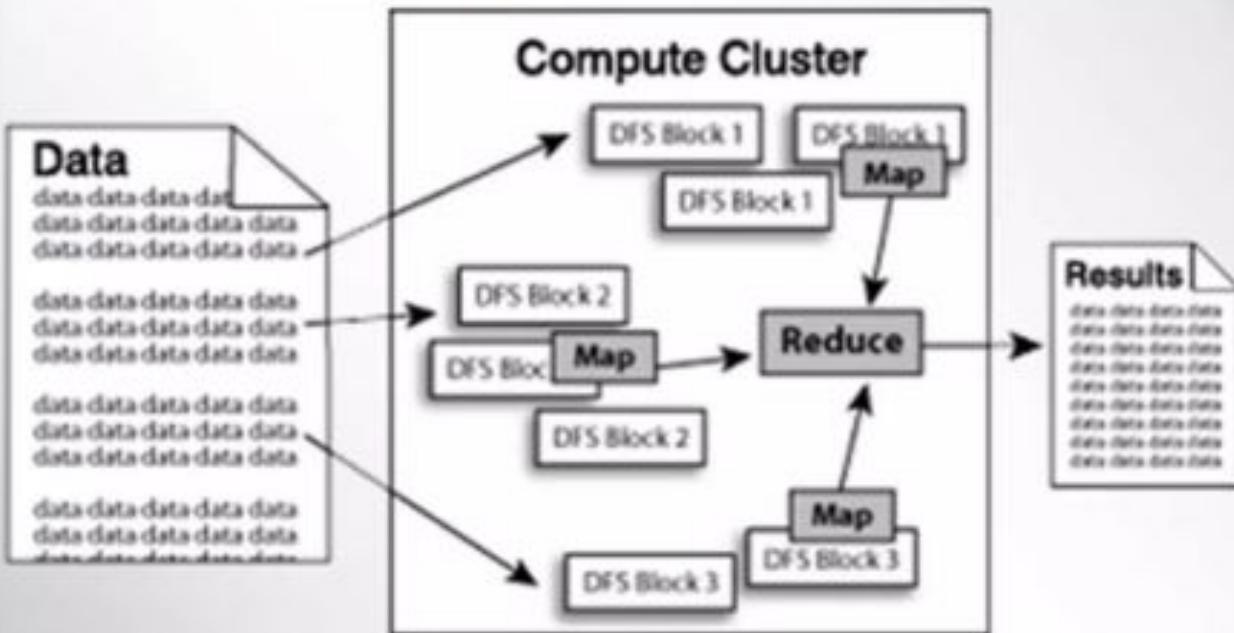
History of Hadoop

- Apache Software Foundation is the developers of Hadoop, and it's co-founders are **Doug Cutting** and **Mike Cafarella**.
- It's co-founder Doug Cutting named it on his son's toy elephant. In October 2003 the first paper release was Google File System.
- In January 2006, MapReduce development started on the Apache Nutch which consisted of around 6000 lines coding for it and around 5000 lines coding for HDFS.
- In April 2006 Hadoop 0.1.0 was released.





hadoop overview



*image courtesy of the
Apache Software Foundation*

Advantages and Disadvantages of Hadoop

Advantages:

- Ability to store a large amount of data.
- High flexibility.
- Cost effective.
- High computational power.
- Tasks are independent.
- Linear scaling.

Disadvantages:

- Not very effective for small data.
- Hard cluster management.
- Has stability issues.
- Security concerns.

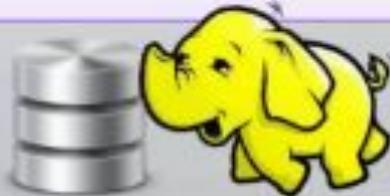
Hadoop Distributed File System

- It has distributed file system known as HDFS and this HDFS splits files into blocks and sends them across various nodes in form of large clusters.
- Also in case of a node failure, the system operates and data transfer takes place between the nodes which are facilitated by HDFS.

Others (For Data Processing)	MapReduce (For Data Processing)
YARN (Resource Management For Cluster)	
HDFS (A Reliable & Redundant Storage)	

Comparing: RDBMS vs. Hadoop

	Traditional RDBMS	Hadoop / MapReduce
Data Size	Gigabytes (Terabytes)	Petabytes (Hexabytes)
Access	Interactive and Batch	Batch – NOT Interactive
Updates	Read / Write many times	Write once, Read many times
Structure	Static Schema	Dynamic Schema
Integrity	High (ACID)	Low
Scaling	Nonlinear	Linear
Query Response Time	Can be near immediate	Has latency (due to batch processing)



Advantages of HDFS:

- It is inexpensive, immutable in nature, stores data reliably, ability to tolerate faults, scalable, block structured, can process a large amount of data simultaneously and many more.

Disadvantages of HDFS:

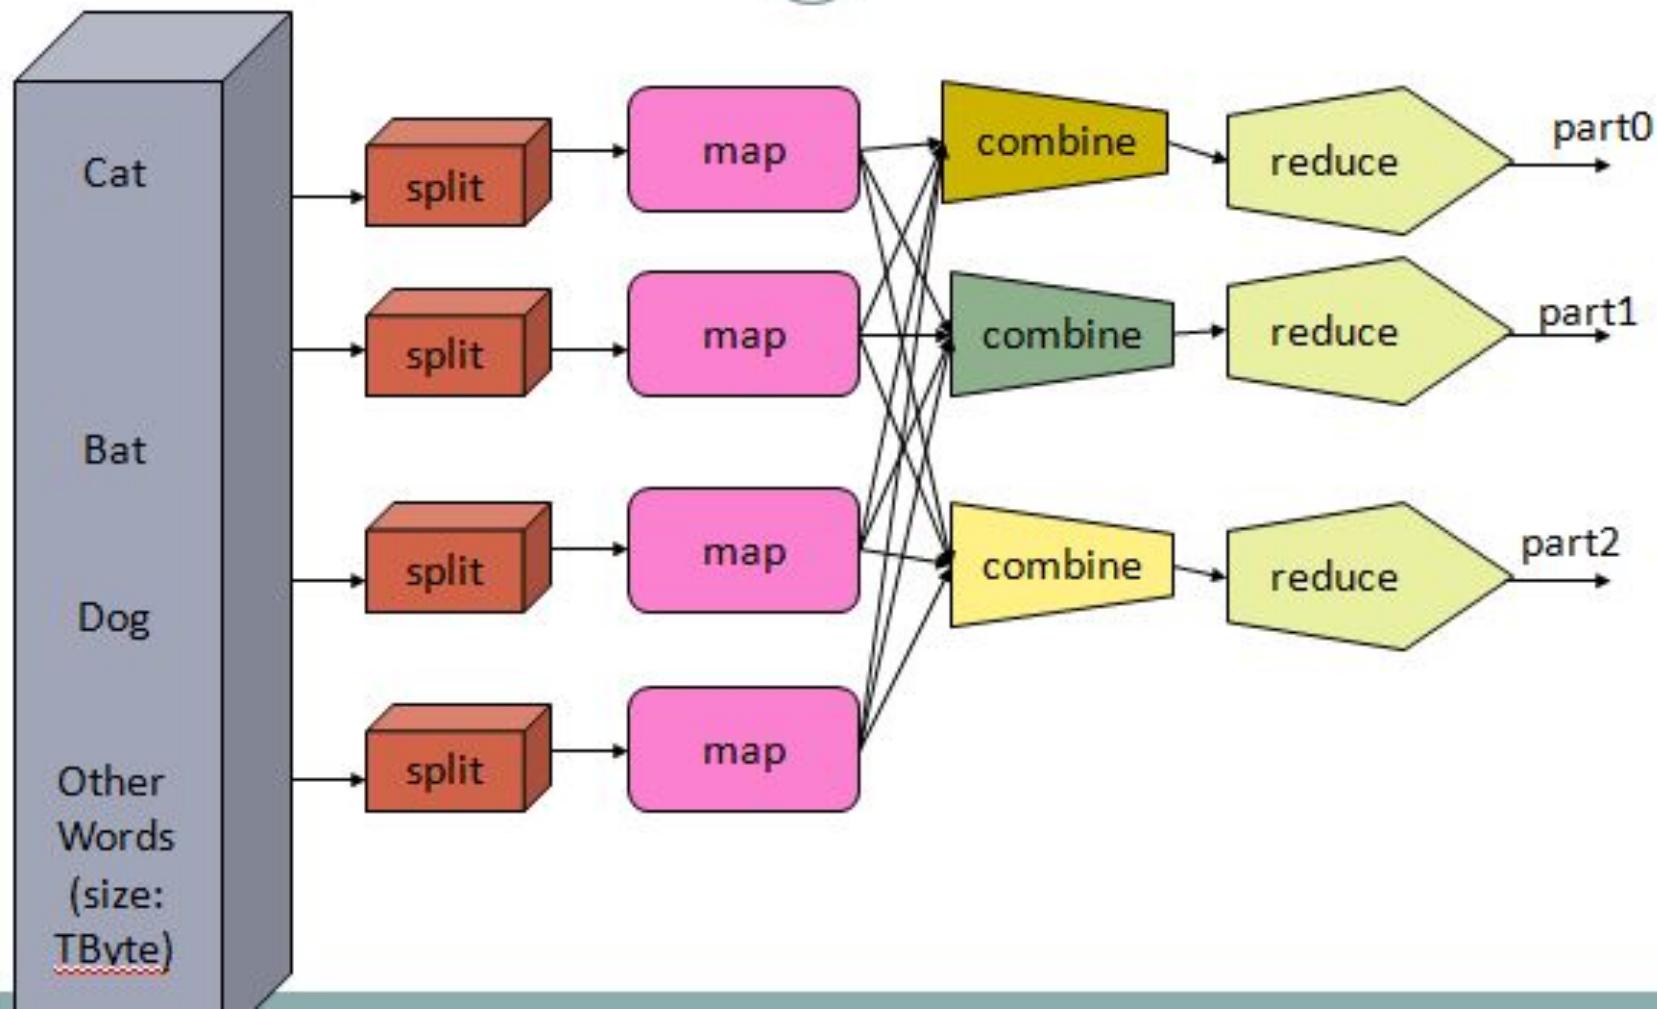
- It's the biggest disadvantage is that it is not fit for small quantities of data. Also, it has issues related to potential stability, restrictive and rough in nature.

Some common frameworks of Hadoop

- **Hive**- It uses HiveQL for data structuring and for writing complicated MapReduce in HDFS.
- **Drill**- It consists of user-defined functions and is used for data exploration.
- **Storm**- It allows real-time processing and streaming of data.
- **Spark**- It contains a Machine Learning Library(MLib) for providing enhanced machine learning and is widely used for data processing. It also supports Java, Python, and Scala.
- **Pig**- It has Pig Latin, a SQL-Like language and performs data transformation of unstructured data.
- **Tez**- It reduces the complexities of Hive and Pig and helps in the running of their codes faster.

Some common frameworks of Hadoop

- **Hive**- It uses HiveQL for data structuring and for writing complicated MapReduce in HDFS.
- **Drill**- It consists of user-defined functions and is used for data exploration.
- **Storm**- It allows real-time processing and streaming of data.
- **Spark**- It contains a Machine Learning Library(MLlib) for providing enhanced machine learning and is widely used for data processing. It also supports Java, Python, and Scala.
- **Pig**- It has Pig Latin, a SQL-Like language and performs data transformation of unstructured data.
- **Tez**- It reduces the complexities of Hive and Pig and helps in the running of their codes faster.

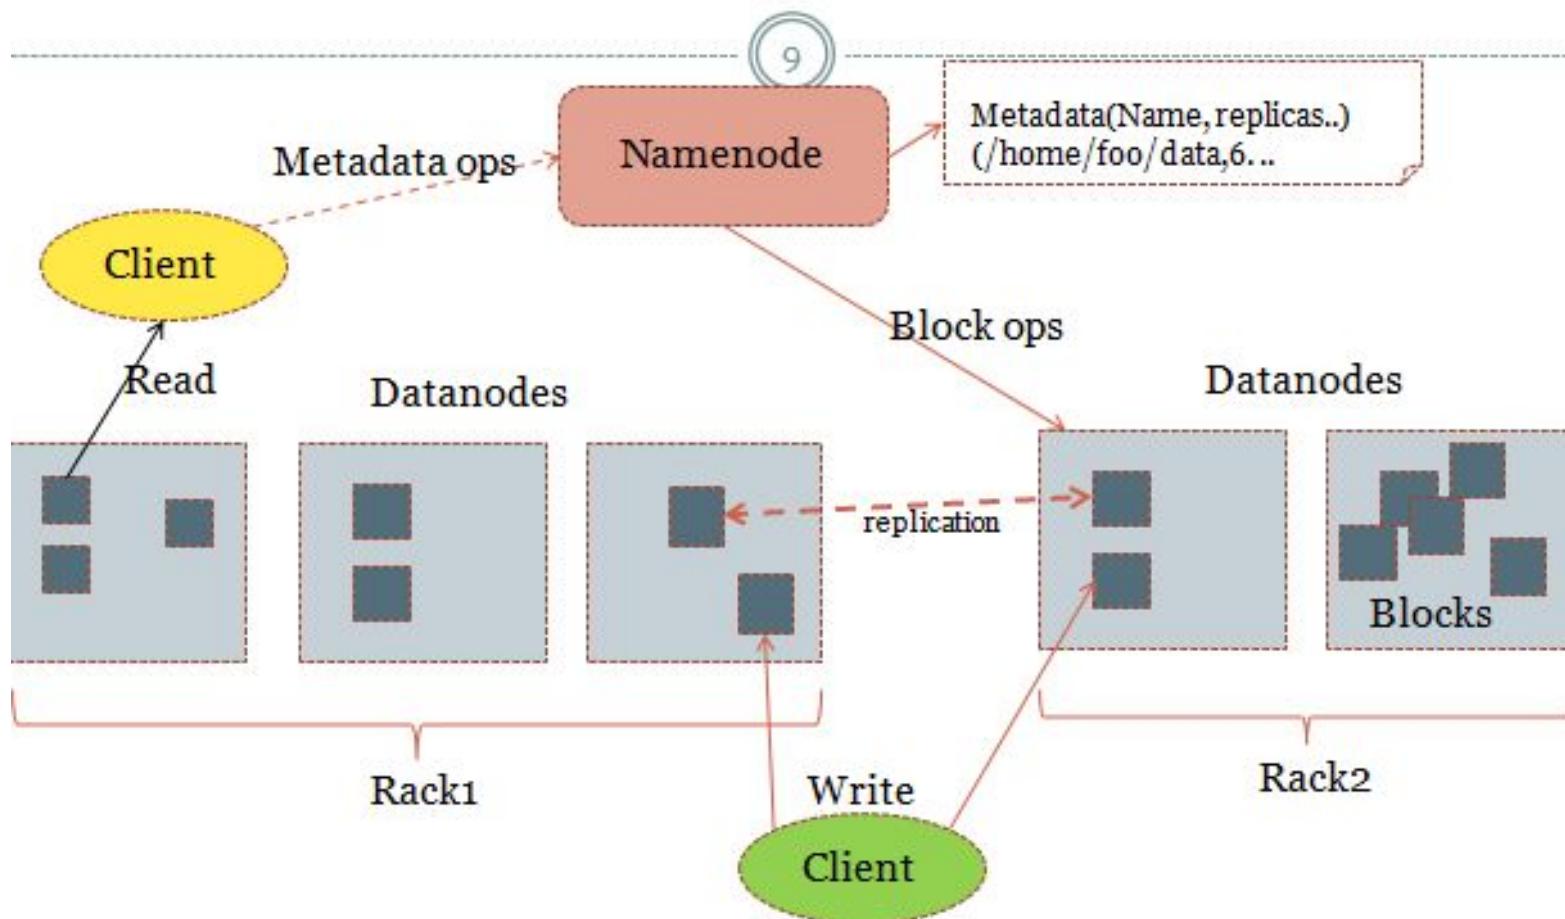


Namenode and Datanodes

8

- Master/slave architecture
- HDFS cluster consists of a single **Namenode**, a master server that manages the file system namespace and regulates access to files by clients.
- There are a number of **DataNodes** usually one per node in a cluster.
- The DataNodes manage storage attached to the nodes that they run on.
- HDFS exposes a file system namespace and allows user data to be stored in files.
- A file is split into one or more blocks and set of blocks are stored in DataNodes.
- DataNodes: serves read, write requests, performs block creation, deletion, and replication upon instruction from Namenode.

HDFS Architecture



Datanode

17

- A Datanode stores data in files in its local file system.
- Datanode has no knowledge about HDFS filesystem
- It stores each block of HDFS data in a separate file.
- Datanode does not create all files in the same directory.
- It uses heuristics to determine optimal number of files per directory and creates directories appropriately:
 - Research issue?
- When the filesystem starts up it generates a list of all HDFS blocks and send this report to Namenode:
Blockreport.

The Communication Protocol

- All HDFS communication protocols are layered on top of the TCP/IP protocol
- A client establishes a connection to a configurable TCP port on the Namenode machine. It talks ClientProtocol with the Namenode.
- The Datanodes talk to the Namenode using Datanode protocol.
- RPC abstraction wraps both ClientProtocol and Datanode protocol.
- Namenode is simply a server and never initiates a request; it only responds to RPC requests issued by DataNodes or clients.

Hadoop Architecture

Hadoop framework includes following four modules:

- ▶ Hadoop Common: These are Java libraries and utilities required by other Hadoop modules. These libraries provides filesystem and OS level abstractions and contains the necessary Java files and scripts required to start Hadoop.
- ▶ Hadoop YARN: This is a framework for job scheduling and cluster resource management.
- ▶ Hadoop Distributed File System (HDFS™): A distributed file system that provides high-throughput access to application data.
- ▶ Hadoop MapReduce: This is YARN-based system for parallel processing of large data sets.

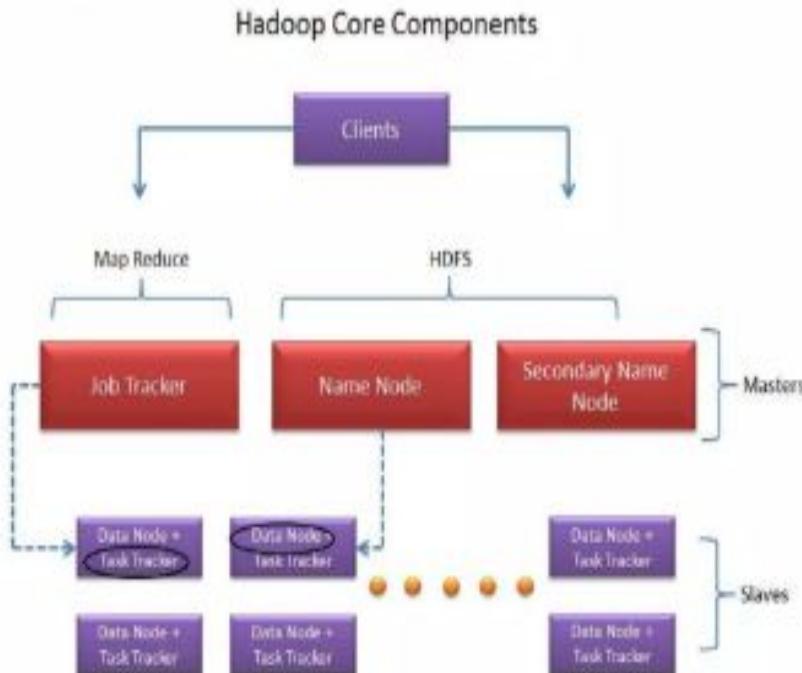
Hadoop Main Components

- ▶ Hadoop consists of MapReduce, the Hadoop distributed file system (HDFS) and a number of related projects such as Apache Hive, HBase and Zookeeper. MapReduce and Hadoop distributed file system (HDFS) are the main component of Hadoop.

Core Components of Hadoop Cluster

Hadoop cluster has 3 components:

- ▶ Client
- ▶ Master
- ▶ Slave



Task Tracker

1. Each Task Tracker is responsible to execute and manage the individual tasks assigned by Job Tracker.
2. Task Tracker also handles the data motion between the map and reduce phases.
3. One Prime responsibility of Task Tracker is to constantly communicate with the Job Tracker the status of the Task.
4. If the JobTracker fails to receive a heartbeat from a TaskTracker within a specified amount of time, it will assume the TaskTracker has crashed and will resubmit the corresponding tasks to other nodes in the cluster.

HDFS & MapReduce

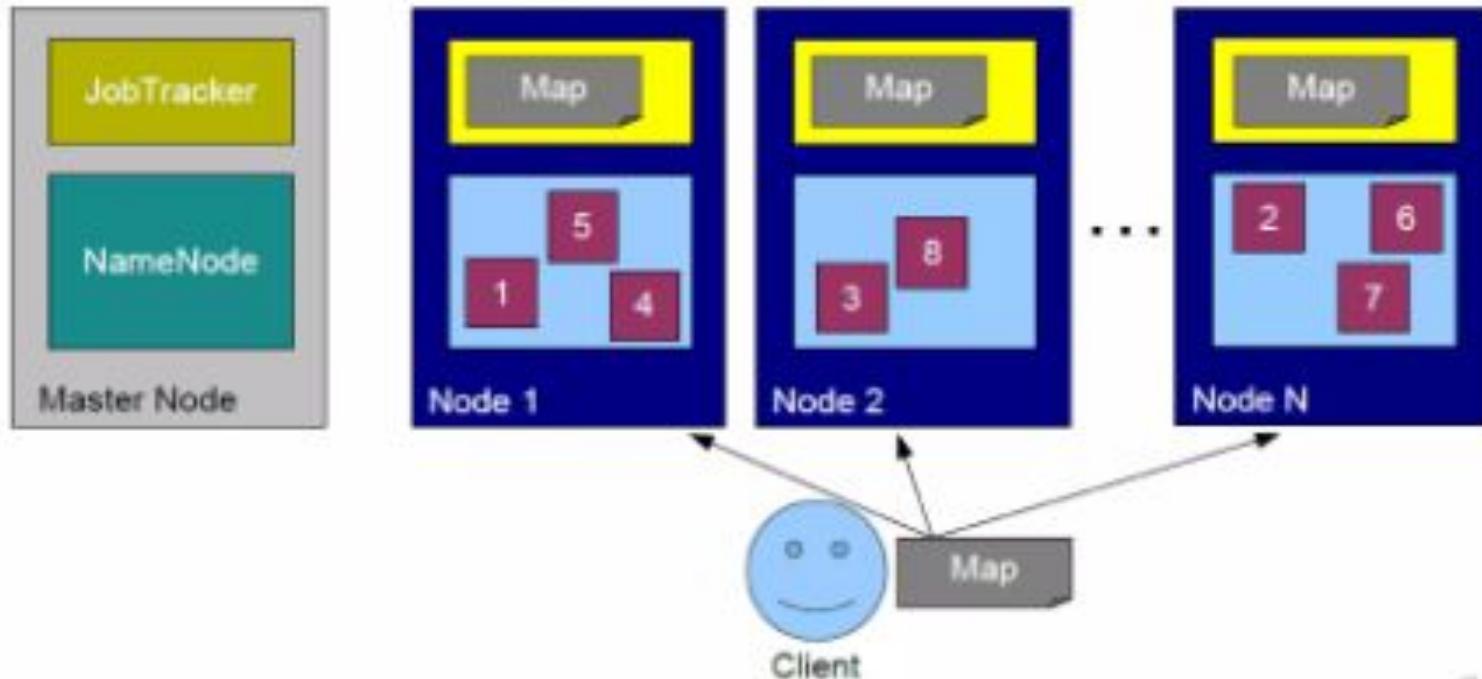
- **Hadoop Distributed File System** 
 - A scalable, Fault tolerant, High performance distributed file system
 - Asynchronous replication
 - Write-once and read-many (WORM)
 - Hadoop cluster with 3 DataNodes minimum
 - Data divided into 64MB (default) or 128MB blocks, each block replicated 3 times (default)
 - No RAID required for DataNode
 - Interfaces: API, Thrift, C Library, FUSE, WebDAV, HTTP, FTP
 - **NameNode** holds filesystem metadata
 - Files are broken up and spread over the **DataNodes**
- **Hadoop Map Reduce** 
 - Software framework for distributed computation
 - Input | Map() | Copy/Sort | Reduce() | Output
 - **JobTracker** schedules and manages jobs
 - **TaskTracker** executes individual map() and reduce() tasks on each cluster node



MapReduce



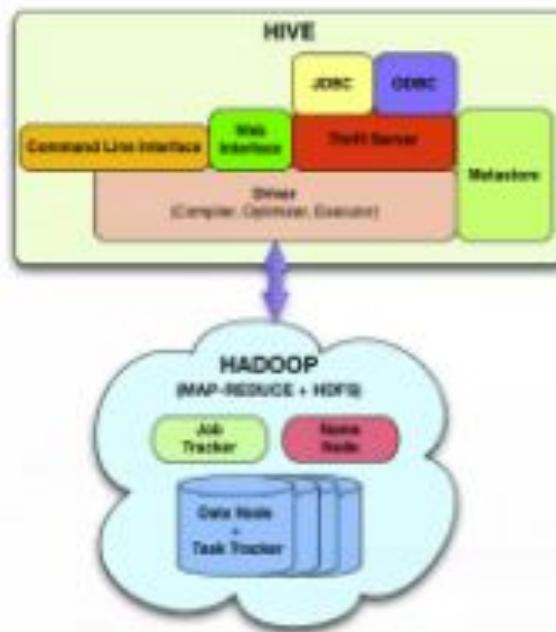
Client's program copied to each node; operates on any blocks present on the node



Hive



- Data Warehouse infrastructure that provides data summarization and ad hoc querying on top of Hadoop
 - MapReduce for execution
 - HDFS for storage
- MetaStore
 - Table/Partitions properties
 - Thrift API : Current clients in Php (Web Interface), Python interface to Hive, Java (Query Engine and CLI)
 - Metadata stored in any SQL backend
- Hive Query Language
 - Basic SQL : Select, From, Join, Group By
 - Equi-Join, Multi-Table Insert, Multi-Group-By
 - Batch query



Pig



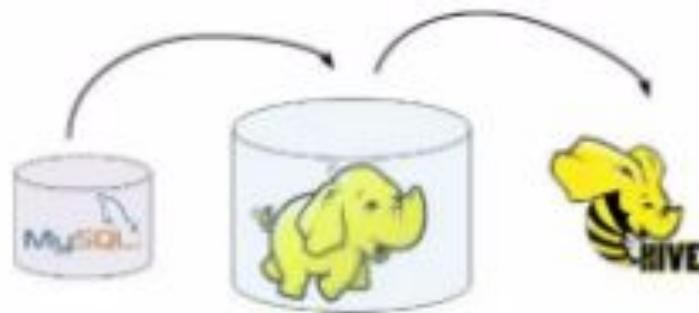
- A high-level data-flow language and execution framework for parallel computation
- Simple to write MapReduce program
- Abstracts you from specific detail
- Focus on data processing
- Data flow
- For data manipulation

PIG Language Example

```
Users = LOAD 'users' AS (name, age);
Fltrd = FILTER Users BY
    age >= 18 AND age <= 25;
Pages = LOAD 'pages' AS (user, url);
Jnd = JOIN Fltrd BY name, Pages BY user;
Grpd = GROUP Jnd BY url;
Smm = FOREACH Grpd GENERATE group,
    COUNT(Jnd) AS clicks;
Srted = ORDER Smm BY clicks DESC;
Top5 = LIMIT Srted 5;
STORE Top5 INTO 'top5sites';
```

Sqoop

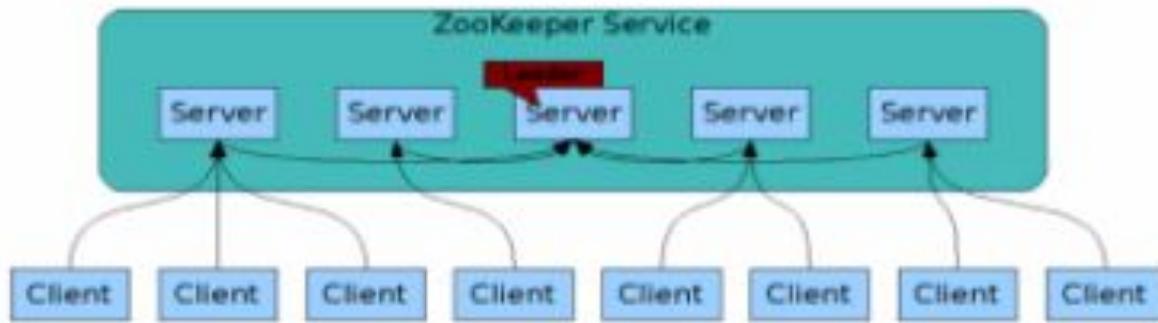
- Sqoop is a tool designed to help users of large data import existing relational databases into their Hadoop clusters
- Automatic data import
- SQL to Hadoop
- Easy import data from many databases to Hadoop
- Generates code for use in MapReduce applications
- Integrates with Hive



Zookeeper



- A high-performance coordination service for distributed applications
- ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services



- All servers store a copy of the data
- A leader is elected at startup
- Followers service clients, all updates go through leader
- Update responses are sent when a majority of servers have persisted the change

Avro



- A data serialization system that provides dynamic integration with scripting languages
- Avro Data
 - Expressive
 - Smaller and Faster
 - Dynamic
 - Schema store with data
 - APIs permit reading and creating
 - Include a file format and a textual encoding
- Avro RPC
 - Leverage versioning support
 - For Hadoop service provide cross-language access



What is big data Analytics?

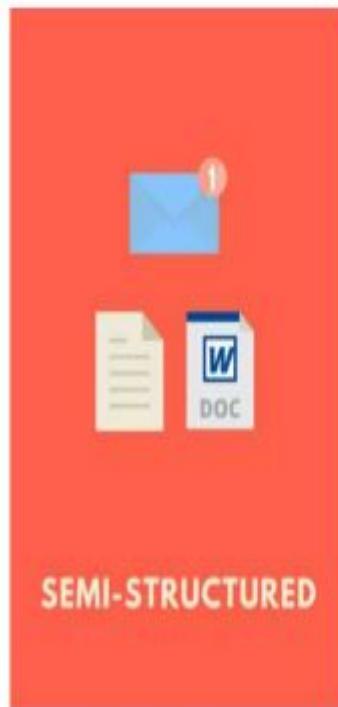
Big data Analytics is a process to extract meaningful insight from big such as hidden patterns, unknown correlations, market trends and customer preferences



Types of Big Data Analytics



Semi-structured Data

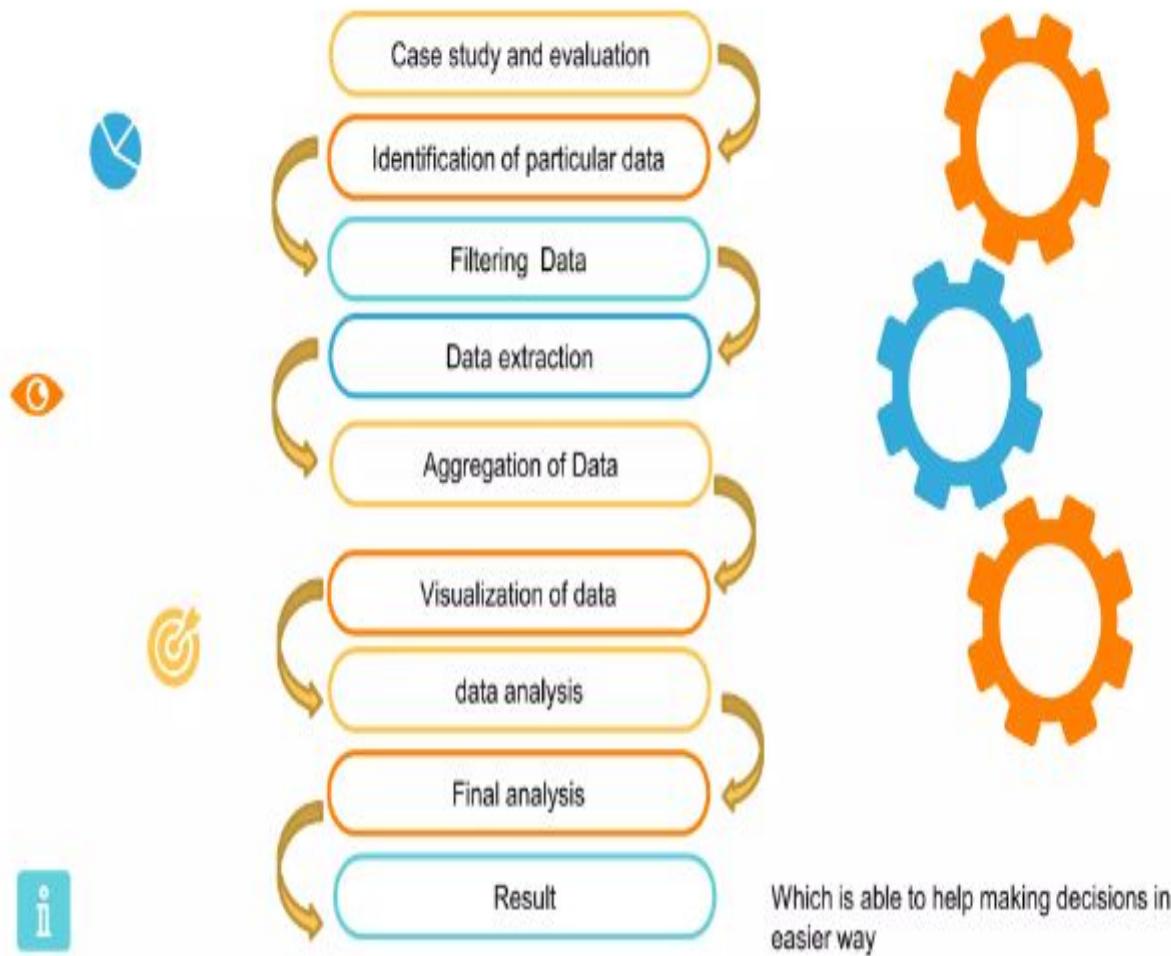


SEMI-STRUCTURED

Semi-Structured Data refers to the data that does not have a proper structure associated with it. For example, the data that is present within the emails, the log files, and the word documents can be referred to as Semi-Structured Data.

```
<rec><name>Prashant Rao</name><sex>Male</sex><age>35</age></rec>
<rec><name>Seema R.</name><sex>Female</sex><age>41</age></rec>
<rec><name>Satish Mane</name><sex>Male</sex><age>29</age></rec>
<rec><name>Subrato Roy</name><sex>Male</sex><age>26</age></rec>
<rec><name>Jeremiah J.</name><sex>Male</sex><age>35</age></rec>
```

Process of Big Data Analytics



Tools used in Big Data Analytics

Tools	MongoDB	MongoDB is used on datasets that change frequently Tools used in big data analytics
	Hadoop	Hadoop helps in storing and analyzing big data Tools used in big data analytics
	Talend	Talend is a tool used for data integration and management Tools used in big data analytics
	Cassandra	It is a distributed database that is used for handling chunks of data
	Storm	It is an open source real time computational system
	Spark	It is used for real time processing and analyzing large amount of data

Benefits Big Data Analytics

- Big data analytics is used for risk management
- Big data analytics is used to improve customer experience
- Big data analytics is used for product development and innovations
- Big data analytics helps in quicker and better decision making in organizations

Google has mastered the domain of big data analytics and it has developed several tools and techniques to capture the data of users which includes their preference, their likes, dislikes, the area of specialization, their requirement etc.



Structured Data



STRUCTURED

Structured Data refers to the data that has a proper structure associated with it. For example, the data that is present within the databases, the CSV files, and the excel spreadsheets can be referred to as Structured Data.

Employee_ID	Employee_Name	Gender	Department	Salary_In_lacs
2365	Rajesh Kulkarni	Male	Finance	650000
3398	Pratibha Joshi	Female	Admin	650000
7465	Shushil Roy	Male	Admin	500000
7500	Shubhojit Das	Male	Finance	500000
7699	Priya Sane	Female	Finance	550000

Unstructured Data



Un-Structured Data refers to the data that does not have any structure associated with it at all. For example, the image files, the audio files, and the video files can be referred to as Un-Structured Data.

Big Data | [Search](#) | [Images](#) | [Maps](#) | [News](#) | [Videos](#) | [Gmail](#) | [Groups](#)

About 3,150,079 results (0.07 seconds)

Big Data & Enterprise - IBM.com
[IBM](#) | [www.ibm.com/big-data-enterprise/](#)
IBM Big Data For Enterprise IBM Big Analytics, Intel & Watson™
Over 20,000 professionals can Google™

100% Optimize for Hadoop - warba.com
[warba.com](#) | [www.warba.com/Hadoop/](#)
We Guarantee the Data Lives in Hadoop 100% reliable real-time availability

Hadoop Big Data - Signaviant.com
[Signaviant](#) | [www.signaviant.com/big-data-training/](#)
Expert Big Data Trainer, Data Flow Project Manager, Cloud Native

Search for hadoop big data

What you missed in Big Data: Hadoop Applications Warba...
Illustrated 20 Steps - 11 mins, 46 sec
Big Data Cloud storage Data-driven apps access documents to the hadoop file system after HDFS's processing that it will become the search

Search for hadoop big data on Google

Big Data Big Data
Rs. 249.00
Amazon.in

Data in Big Data
Rs. 499.00
Amazon.in

Big Data - The Definitive Guide
Rs. 403.00
Amazon.in

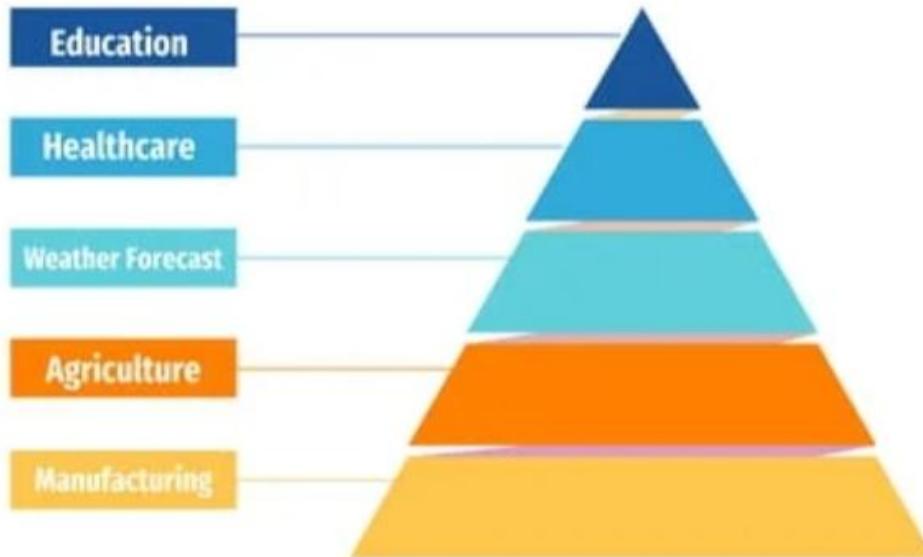
Hadoop Big Data
Rs. 395.00
Amazon.in

Hadoop Big Data
Rs. 249.00
Amazon.in

Big Data Simplified
Rs. 1,799.00
Amazon.in

Hadoop The Definitive Guide
Rs. 393.00
Amazon.in

Big Data application domains



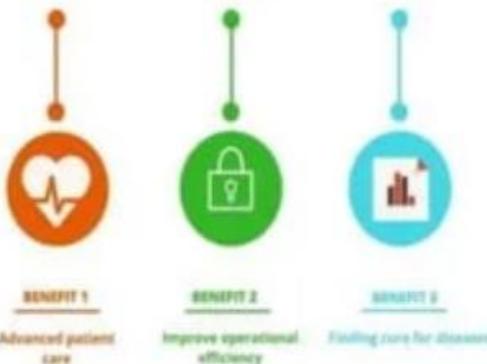
Education



- Enhancing Student Results
- A better Grading System
- Gaining Attention
- Customized Programs
- Reducing The Number of Dropouts

HealthCare

BENEFITS OF BIG DATA IN HEALTHCARE - A REVOLUTION IN THE MAKING



Health Tracking

Improve the care delivery system / machinery

Fraud detection and prevention

Real-time alerts

Weather Forecast



- Estimates of areas where flooding is likely to be most severe
- The strength and direction of tropical storms
- The most likely amount of snow or rain that will fall in a specific area
- The most likely locations of downed power lines
- Estimates of areas where wind speeds are likely to be greatest
- The locations where bridges and roads most likely to be damaged by storms
- The likelihood of flights being cancelled at specific airports

Agriculture



- Boosting productivity
- Access to plant genome information
- Predicting yields
- Risk management
- Food safety
- Savings

Manufacturing



Product quality

Defects tracking

Supply planning

Manufacturing process defect tracking

Testing and simulation of new manufacturing processes

Google Big Data

Google developed several open source tools and techniques that are extensively used in big data ecosystem. With the help of different big data tools and techniques, Google is now capable of exploring millions of websites and **fetch you the right answer or information within milliseconds.**



Google Big Data

The first question that comes to our mind is how can Google perform such complex operations so efficiently?

The answer is Big data analytics.

Big Data tools and techniques to understand our requirements based on several parameters like search history, locations, trends etc.

Google effortlessly displays the complex calculations which are designed to match the user's requirement.



Google always wanted to develop a search engine that has the ability to think like a human and understand the phrase, logic, and goal of any search query. Semantics has helped Google to accomplish this task to look beyond the literal meaning of any phrase of a search query.



Google Has adopted the following techniques using Big Data Analytics



IBM's Weather Forecasting

One example of an application of big data to weather forecasting is IBM's Deep Thunder. Unlike many weather forecasting systems, which give general information about a broad geographical region, Deep Thunder provides forecasts for extremely specific locations, such as a single airport, so that local authorities can get critically important information in real time.





New York Stock Exchange

The New York Stock Exchange generates about one terabyte of new trade data per day.

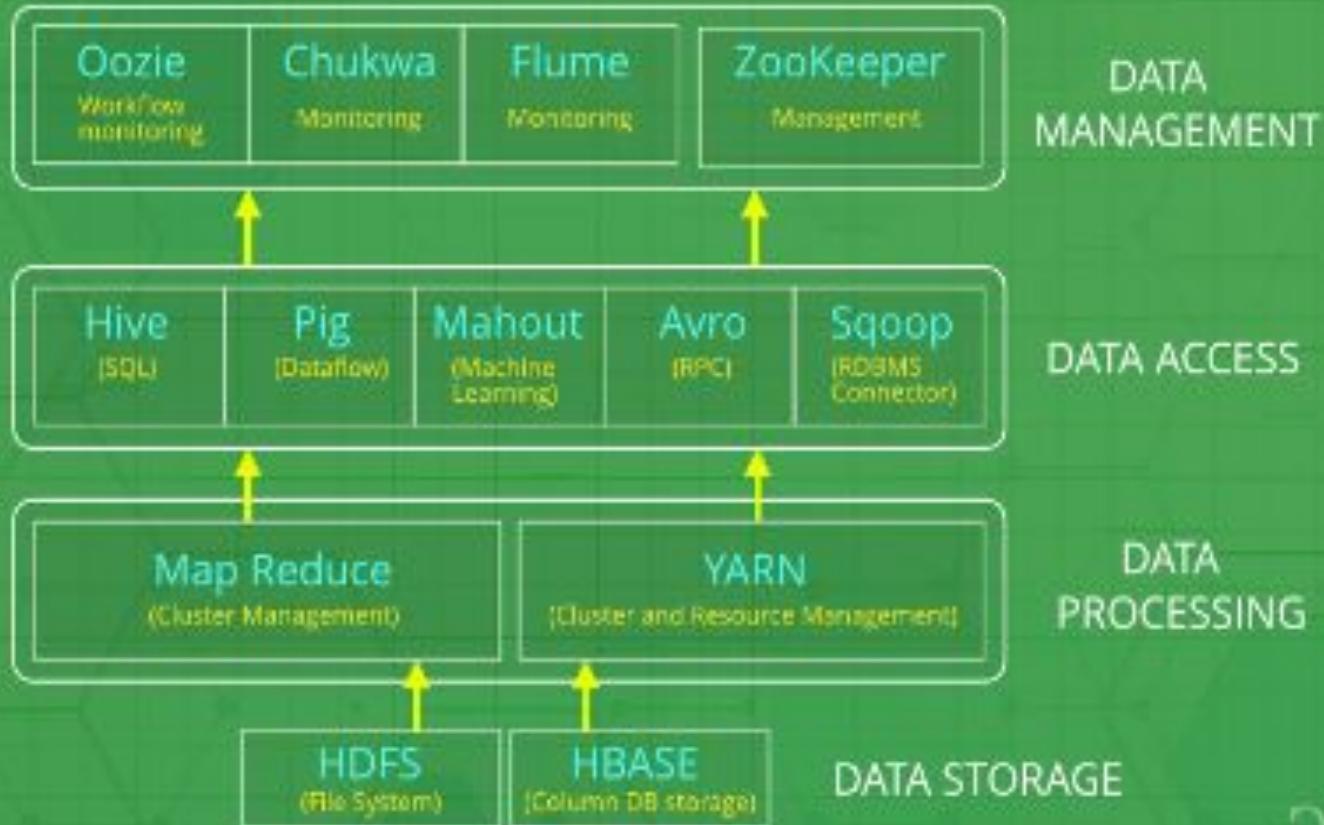
Social Media

500TB+ data is generated everyday only on Facebook. This data is mainly generated in terms of photo and video uploads, message exchanges, putting comments etc.

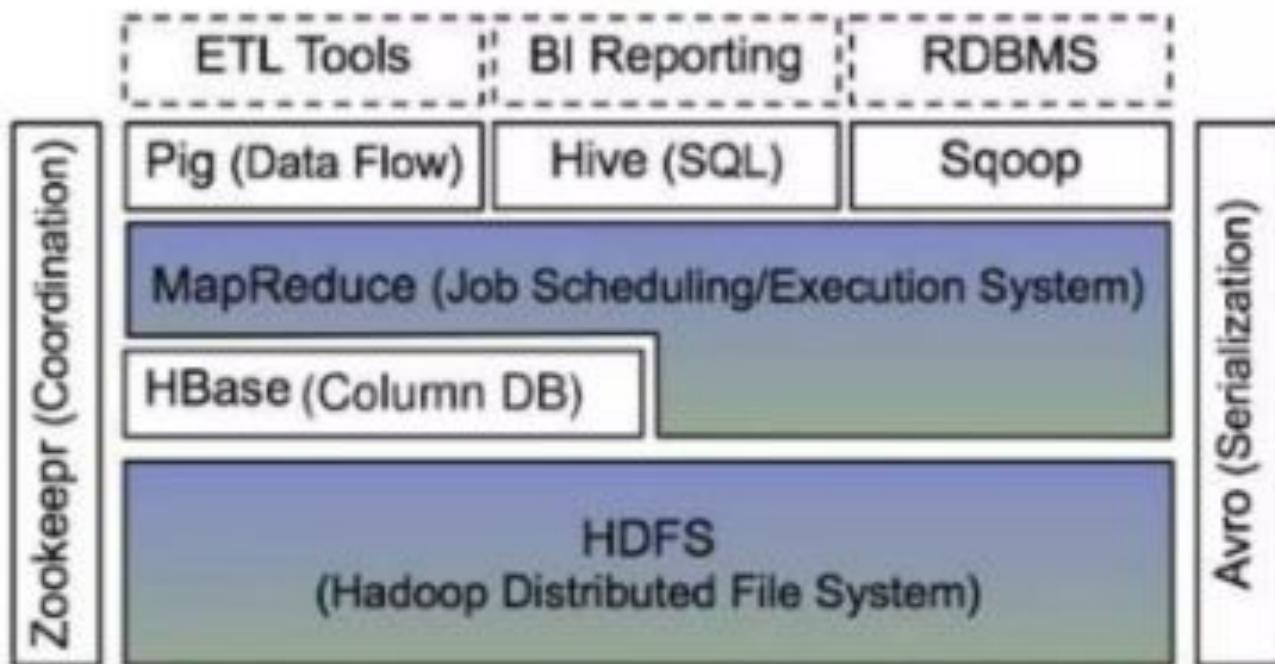
Air Industry

A single Jet engine can generate 10+terabytes of data in 30 minutes of flight time. With many thousand flights per day, generation of data reaches up to many Petabytes.

Hadoop Ecosystem

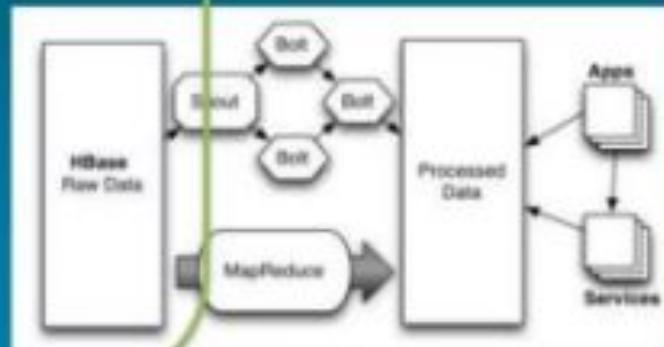


The Hadoop Ecosystem

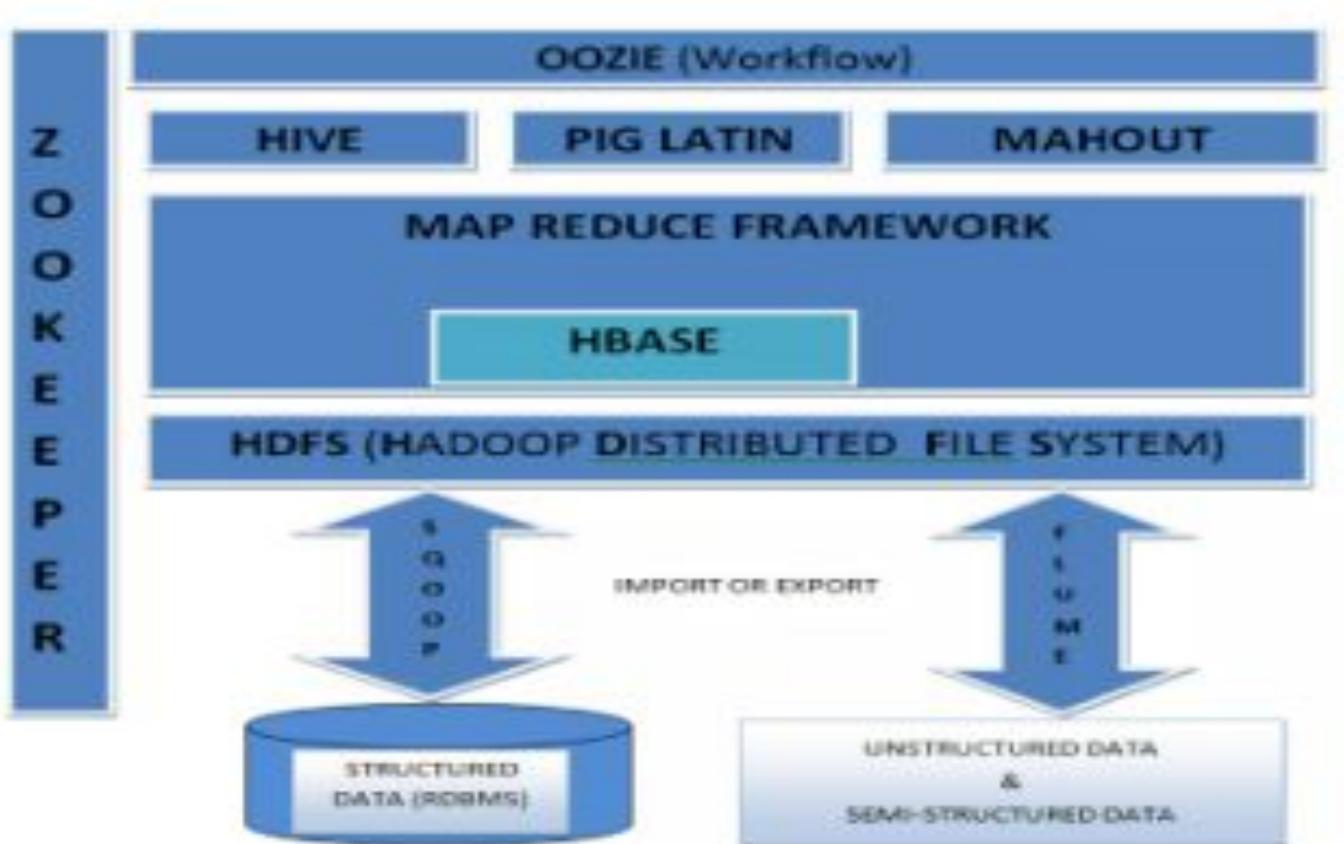


Roles in Hadoop Ecosystem

- MapReduce for indexing, learning
- HBase for storage and fast access
- Also: Storm for incremental update
- And: relational DB for most recent derived data
- API façade for input; API for querying learning



Working Of Ecosystem



Introduction: *Hadoop Ecosystem* is a platform or a suite which provides various services to solve the big data problems. It includes Apache projects and various commercial tools and solutions. There are *four major elements of Hadoop* i.e. **HDFS, MapReduce, YARN, and Hadoop Common**. Most of the tools or solutions are used to supplement or support these major elements. All these tools work collectively to provide services such as absorption, analysis, storage and maintenance of data etc.

Following are the components that collectively form a Hadoop ecosystem:

Following are the components that collectively form a Hadoop ecosystem:

- **HDFS**: Hadoop Distributed File System
- **YARN**: Yet Another Resource Negotiator
- **MapReduce**: Programming based Data Processing
- **Spark**: In-Memory data processing
- **PIG, HIVE**: Query based processing of data services
- **HBase**: NoSQL Database
- **Mahout, Spark MLLib**: [Machine Learning](#) algorithm libraries
- **Solar, Lucene**: Searching and Indexing
- **Zookeeper**: Managing cluster
- **Oozie**: Job Scheduling

Apache HBase:

- It's a NoSQL database which supports all kinds of data and thus capable of handling anything of Hadoop Database. It provides capabilities of Google's BigTable, thus able to work on Big Data sets effectively.
- At times where we need to search or retrieve the occurrences of something small in a huge database, the request must be processed within a short quick span of time. At such times, HBase comes handy as it gives us a tolerant way of storing limited data

HDFS:

- HDFS is the primary or major component of Hadoop ecosystem and is responsible for storing large data sets of structured or unstructured data across various nodes and thereby maintaining the metadata in the form of log files.
- HDFS consists of two core components i.e.
 1. Name node
 2. Data Node
- Name Node is the prime node which contains metadata (data about data) requiring comparatively fewer resources than the data nodes that stores the actual data. These data nodes are commodity hardware in the distributed environment. Undoubtedly, making Hadoop cost effective.
- HDFS maintains all the coordination between the clusters and hardware, thus working at the heart of the system.

YARN:

- Yet Another Resource Negotiator, as the name implies, YARN is the one who helps to manage the resources across the clusters. In short, it performs scheduling and resource allocation for the Hadoop System.
- Consists of three major components i.e.
 1. Resource Manager
 2. Nodes Manager
 3. Application Manager
- Resource manager has the privilege of allocating resources for the applications in a system whereas Node managers work on the allocation of resources such as CPU, memory, bandwidth per machine and later on acknowledges the resource manager. Application manager works as an interface between the resource manager and node manager and performs negotiations as per the requirement of the two.

YARN: Taking Hadoop Beyond Batch

Store ALL DATA in one place...

Interact with that data in MULTIPLE WAYS

with Predictable Performance and Quality of Service



Hadoop 1 Limitations

Scalability

Max Cluster size ~5,000 nodes

Max concurrent tasks ~40,000

Coarse Synchronization in JobTracker

Availability

Failure Kills Queued & Running Jobs

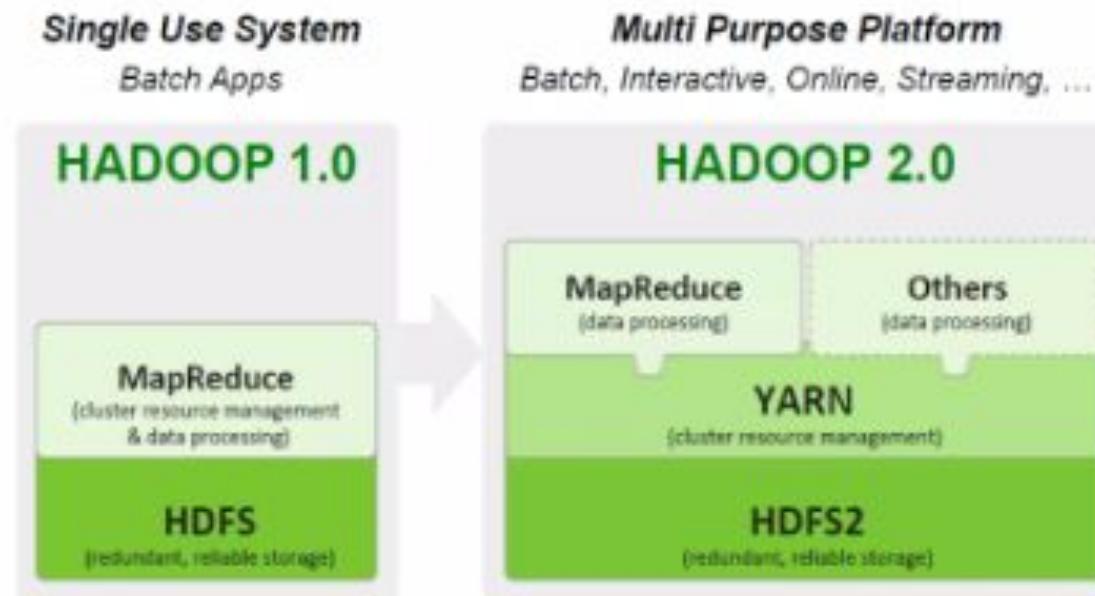
Hard partition of resources into map and reduce slots

Non-optimal Resource Utilization

Lacks Support for Alternate Paradigms and Services

Iterative applications in MapReduce are 10x slower

Our Vision: Hadoop as Next-Gen Platform



MapReduce:

- By making the use of distributed and parallel algorithms, MapReduce makes it possible to carry over the processing's logic and helps to write applications which transform big data sets into a manageable one.
- MapReduce makes the use of two functions i.e. Map() and Reduce() whose task is:
 1. *Map()* performs sorting and filtering of data and thereby organizing them in the form of group. Map generates a key-value pair based result which is later on processed by the Reduce() method.
 2. *Reduce()*, as the name suggests does the summarization by aggregating the mapped data. In simple, Reduce() takes the output generated by Map() as input and combines those tuples into smaller set of tuples.

PIG:

Pig was basically developed by Yahoo which works on a pig Latin language, which is Query based language similar to SQL.

- It is a platform for structuring the data flow, processing and analyzing huge data sets.
- Pig does the work of executing commands and in the background, all the activities of MapReduce are taken care of. After the processing, pig stores the result in HDFS.
- Pig Latin language is specially designed for this framework which runs on Pig Runtime. Just the way Java runs on the JVM.
- Pig helps to achieve ease of programming and optimization and hence is a major segment of the Hadoop Ecosystem.

HIVE:

- With the help of SQL methodology and interface, HIVE performs reading and writing of large data sets. However, its query language is called as HQL (Hive Query Language).
- It is highly scalable as it allows real-time processing and batch processing both. Also, all the SQL datatypes are supported by Hive thus, making the query processing easier.
- Similar to the Query Processing frameworks, HIVE too comes with two components: *JDBC Drivers* and *HIVE Command Line*.
- JDBC, along with ODBC drivers work on establishing the data storage permissions and connection whereas HIVE Command line helps in the processing of queries.

Apache Spark:

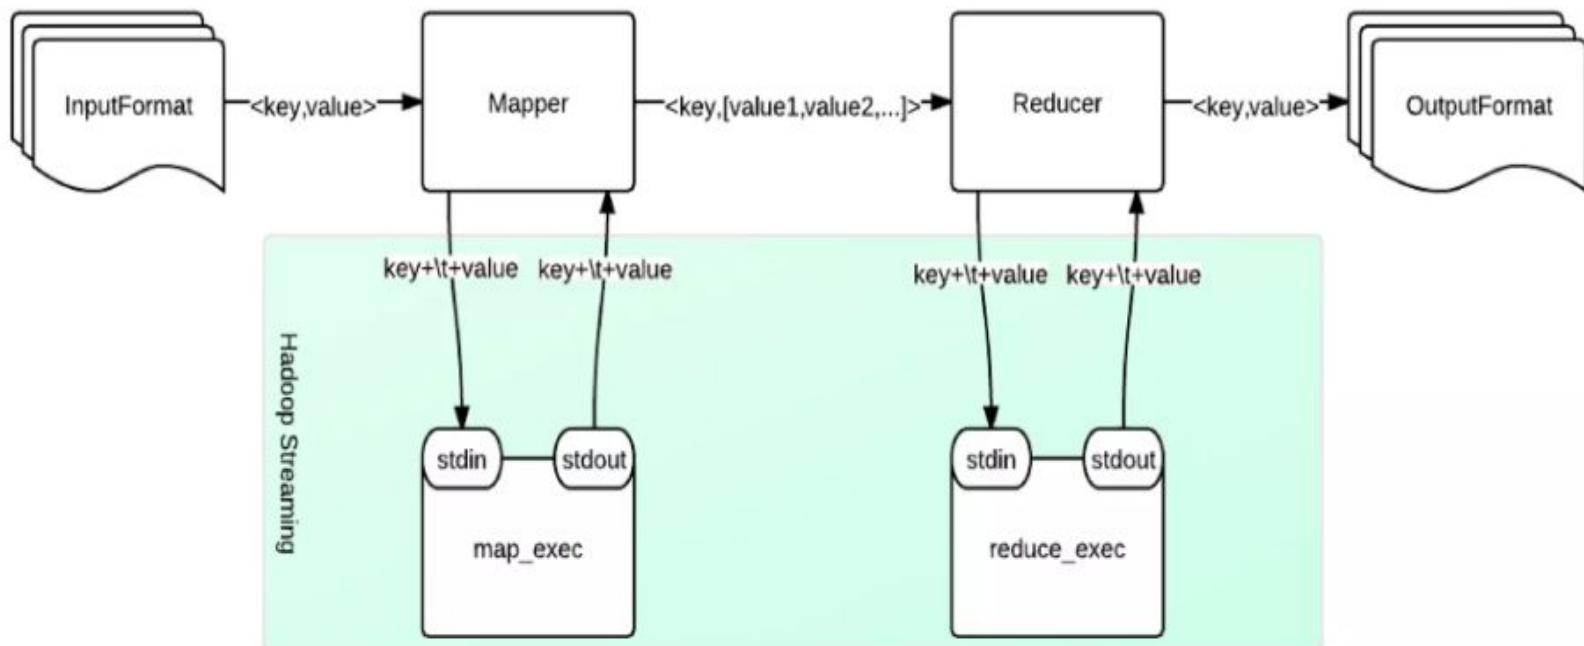
- It's a platform that handles all the process consumptive tasks like batch processing, interactive or iterative real-time processing, graph conversions, and visualization, etc.
- It consumes in memory resources hence, thus being faster than the prior in terms of optimization.
- Spark is best suited for real-time data whereas Hadoop is best suited for structured data or batch processing, hence both are used in most of the companies interchangeably.

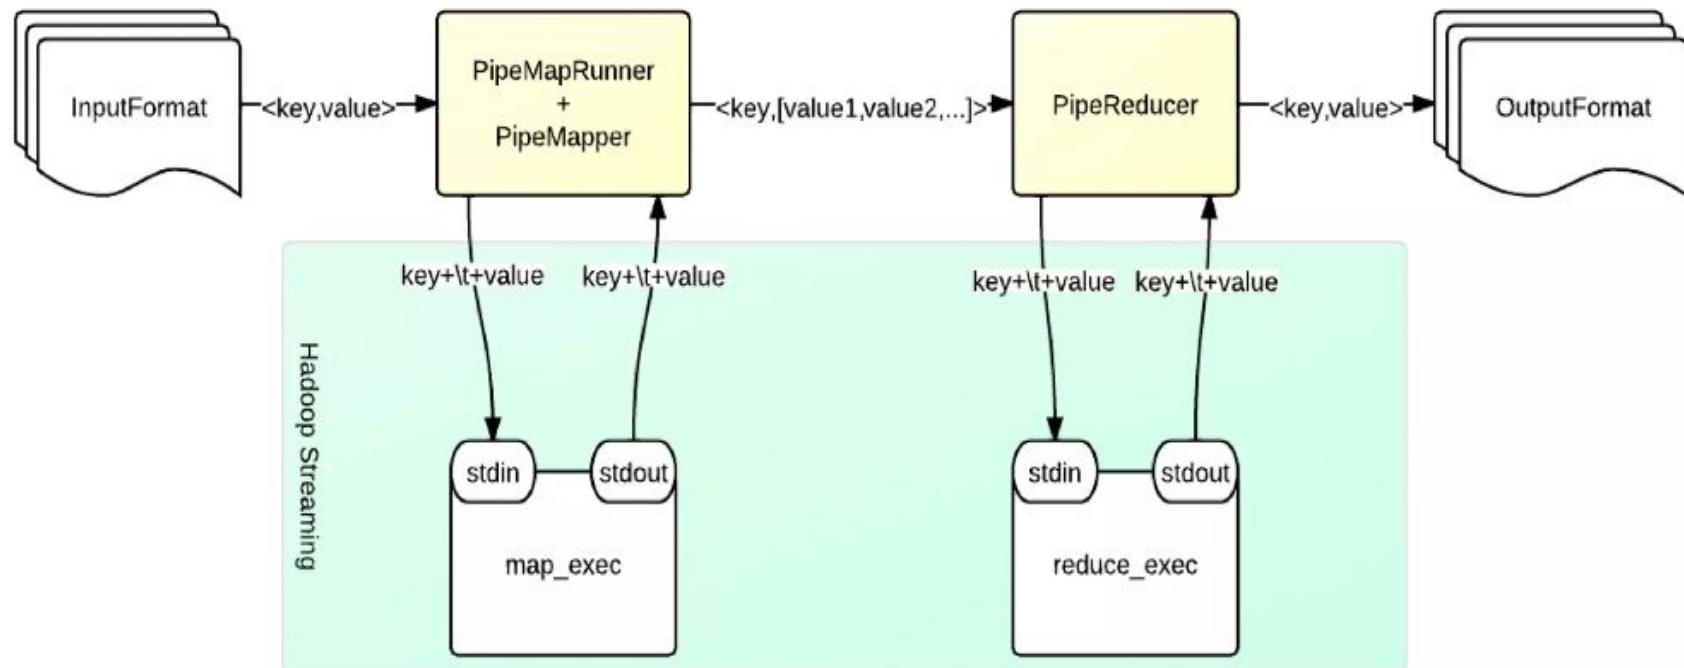
Apache Spark:

- It's a platform that handles all the process consumptive tasks like batch processing, interactive or iterative real-time processing, graph conversions, and visualization, etc.
- It consumes in memory resources hence, thus being faster than the prior in terms of optimization.
- Spark is best suited for real-time data whereas Hadoop is best suited for structured data or batch processing, hence both are used in most of the companies interchangeably.

- **Zookeeper:** There was a huge issue of management of coordination and synchronization among the resources or the components of Hadoop which resulted in inconsistency, often. Zookeeper overcame all the problems by performing synchronization, inter-component based communication, grouping, and maintenance.
- **Oozie:** Oozie simply performs the task of a scheduler, thus scheduling jobs and binding them together as a single unit. There are two kinds of jobs .i.e Oozie workflow and Oozie coordinator jobs. Oozie workflow is the jobs that need to be executed in a sequentially ordered manner whereas Oozie Coordinator jobs are those that are triggered when some data or external stimulus is given to it.

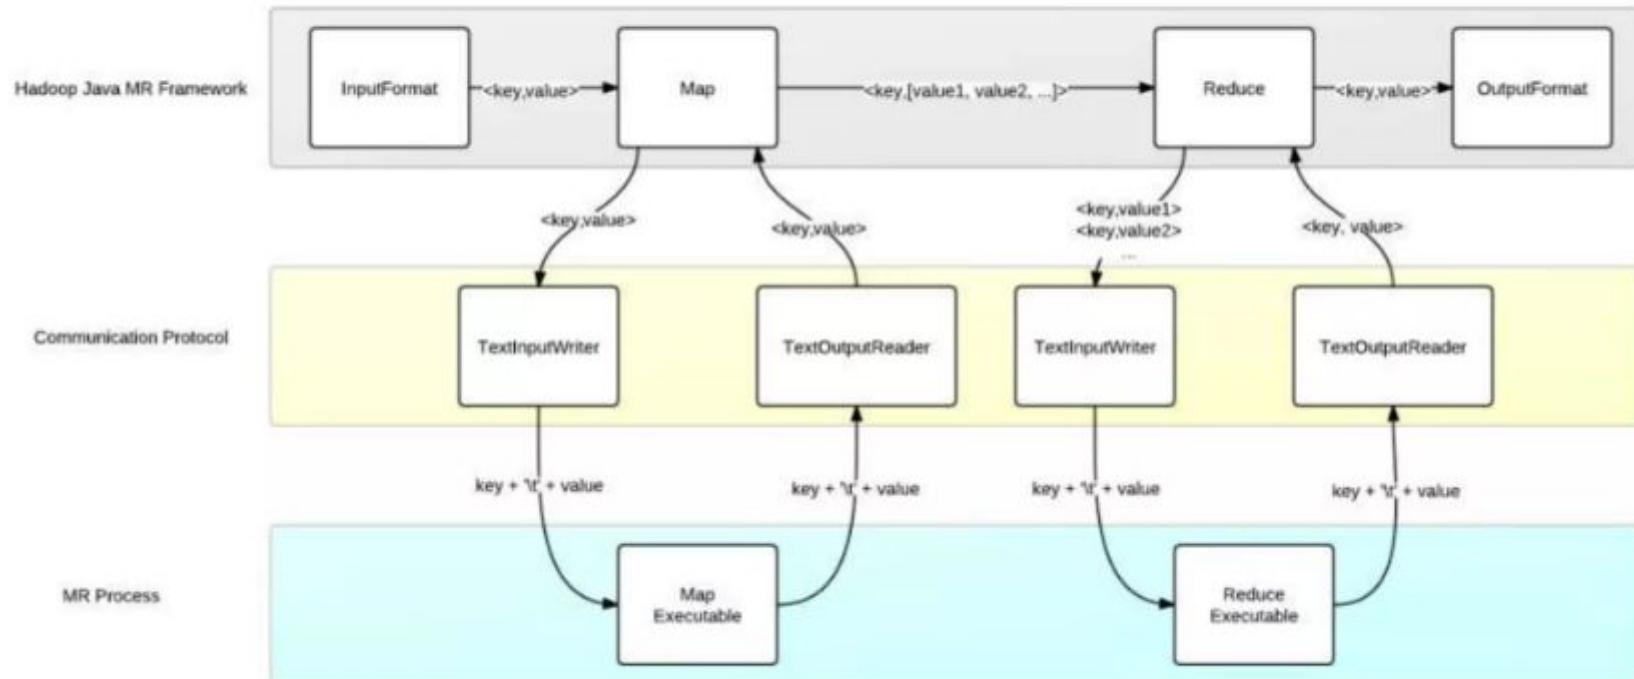
HADOOP PIPES





Streaming Data Flow

HANBORQ



Streaming I/O Format

- `-inputFormat <javaClassName>`
 - `JobConf.setInputFormat()`
- `-outputFormat <javaClassName>`

Hadoop Streaming :

22

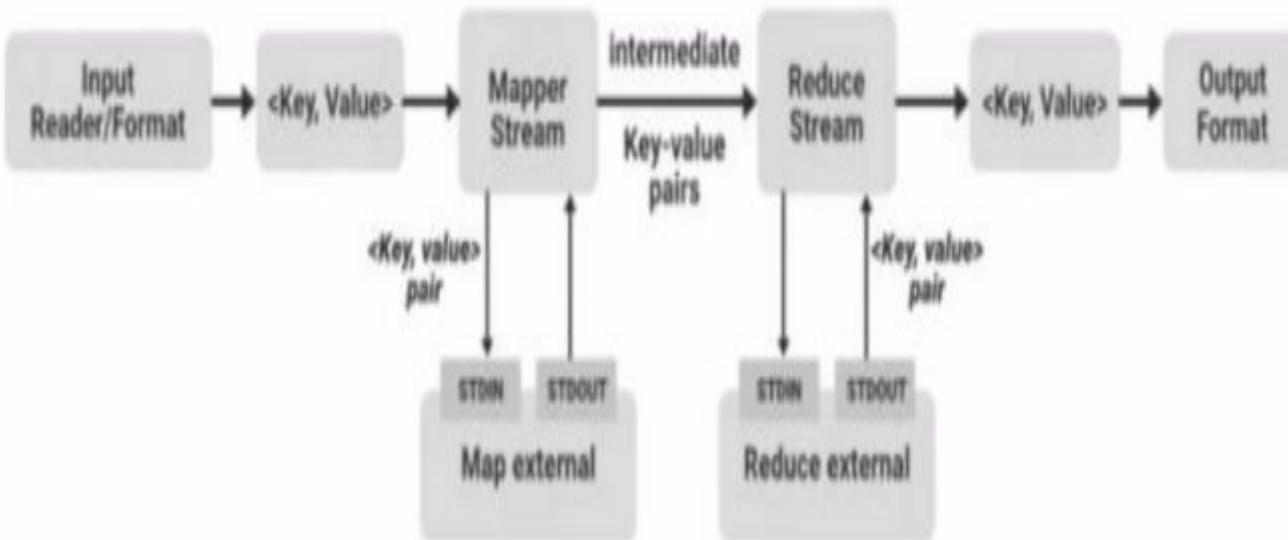
- It is a feature that comes with a Hadoop distribution that allows developers or programmers to write the Map-Reduce program using different programming languages like Ruby, Perl, Python, C++, etc.
- We can use any language that can read from the standard input(STDIN) like keyboard input and all and write using standard output(STDOUT).
- Although Hadoop Framework is completely written in java programs for Hadoop do not necessarily need to code in Java programming language.

Hadoop Streaming :

23

- The main advantage of the Hadoop streaming utility is that it allows Java as well as non-Java programmed MapReduce jobs to be executed over Hadoop clusters.
- The Hadoop streaming supports the Perl, Python, PHP, R, and C++ Translate the application logic into the Mapper and Reducer sections with the key and value output elements.
- Three main components: Mapper, Reducer, and Driver.

Hadoop Streaming

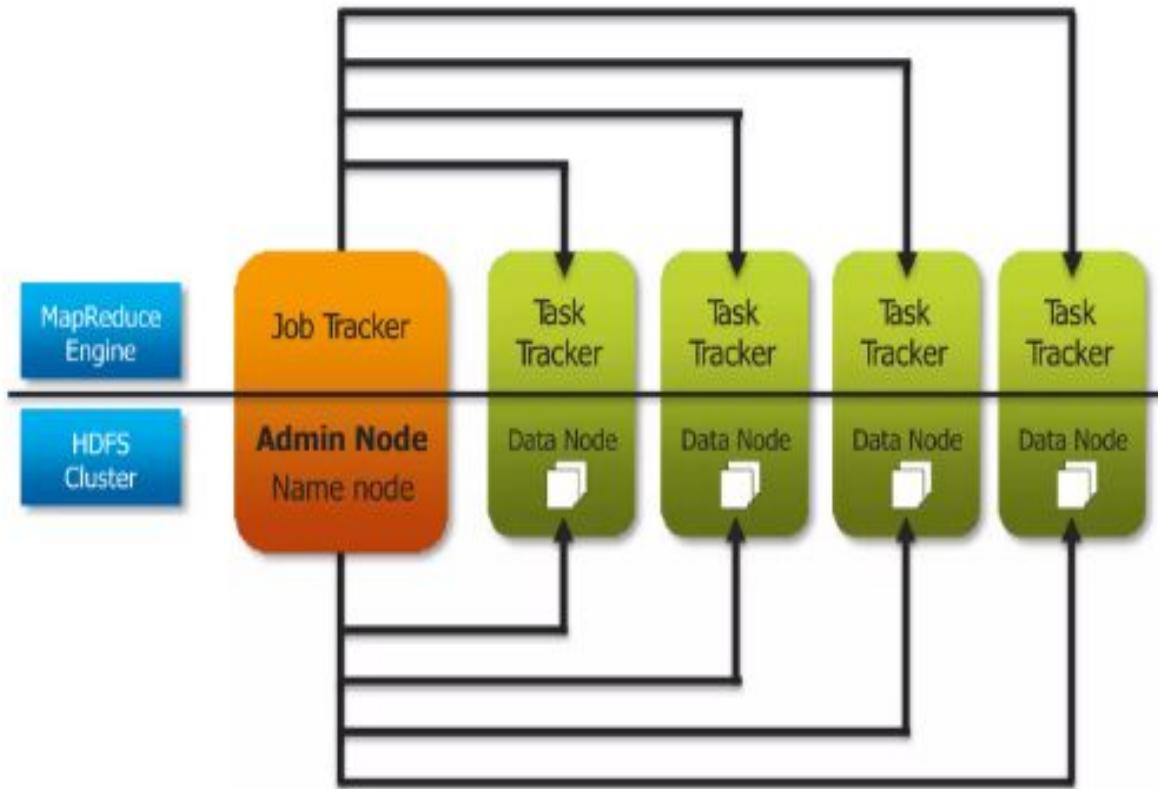


MAP REDUCE

Map Reduce

- HDFS handles the Distributed Filesystem layer
- MapReduce is a programming model for data processing.
- MapReduce
 - Framework for parallel computing
 - Programmers get simple API
 - Don't have to worry about handling
 - parallelization
 - data distribution
 - load balancing
 - fault tolerance
 - Allows one to process huge amounts of data (terabytes and petabytes) on thousands of processors

Map Reduce Concepts (Hadoop-1.0)



Map Reduce Concepts

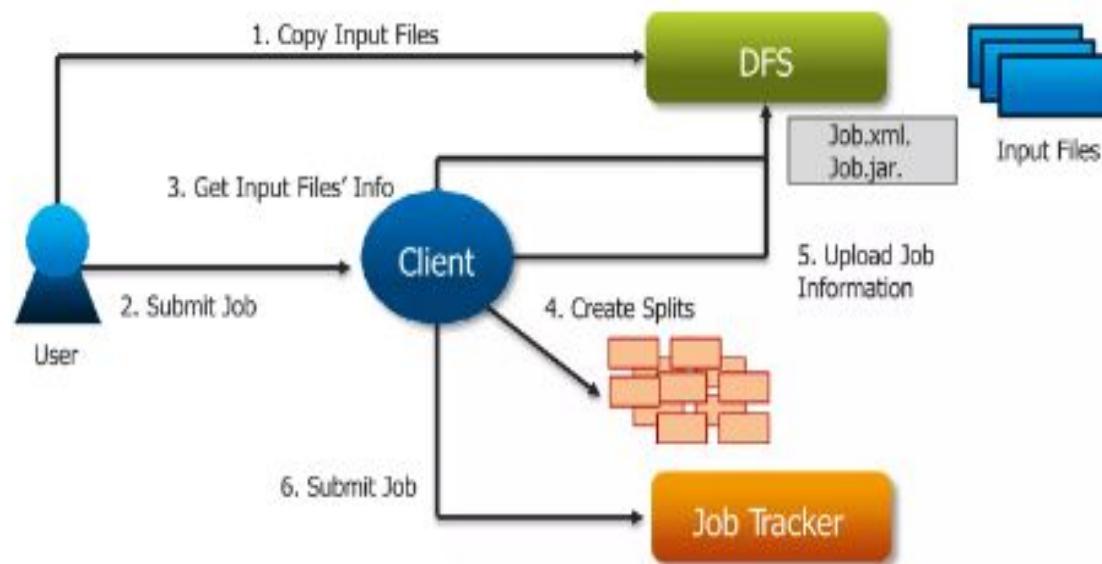
- ❖ **Job Tracker**

- The Job-Tracker is responsible for accepting jobs from clients, dividing those jobs into tasks, and assigning those tasks to be executed by worker nodes.

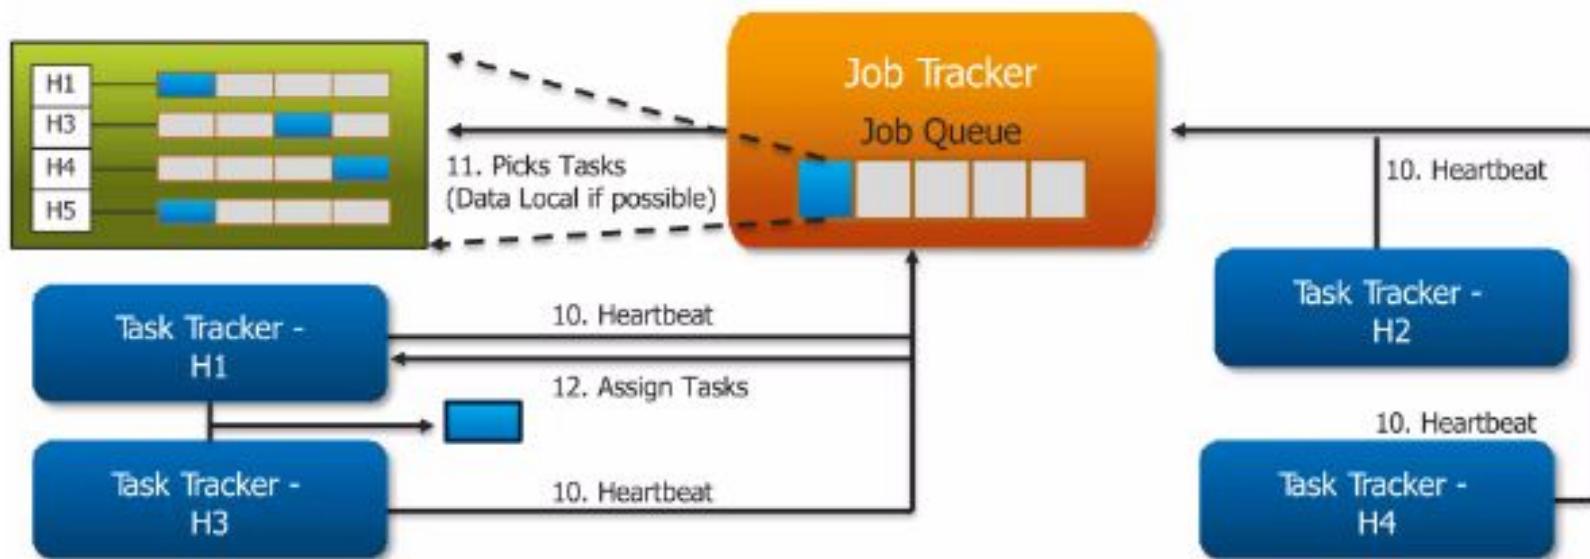
- ❖ **Task Tracker**

- Task-Tracker is a process that manages the execution of the tasks currently assigned to that node. Each Task Tracker has a fixed number of slots for executing tasks (two maps and two reduces by default).

Job Tracker



Job Tracker



Understanding Data Transformations

- In order to write MapReduce applications you need to have an understanding of how data is transformed as it executes in the MapReduce framework.

From start to finish, there are four fundamental transformations. Data is:

- Transformed from the input files and fed into the mappers
- Transformed by the mappers
- Sorted, merged, and presented to the reducer
- Transform by reducers and written to output files

Solving a Programming Problem using MapReduce



Example Data Set

2000	2025191	1788950	236241	1544607	1458185	86422	480584	330765	149819
2001	1991082	1862846	128236	1483563	1516008	-32445	507519	346838	160681
2002	1853136	2010894	-157758	1337815	1655232	-317417	515321	355662	159659
2003	1782314	2159899	-377585	1258472	1796890	-538418	523842	363009	160833
2004	1880114	2292841	-812727	1345369	1913330	-567961	534745	379511	155234
2005	2153611	2471957	-318346	1576135	2069746	-493611	577476	402211	175265
2006	2406869	2655050	-248181	1798487	2232981	-434494	608382	422069	186313
2007	2567985	2728686	-360701	1932896	2275049	-342153	635089	453637	181452
2008	2523991	2982544	-458553	1865945	2507793	-641848	658046	474751	183295
2009	2104989	3517677	-1812688	1450980	3000661	-1549681	654009	517016	136993
2010	2162706	3457079	-1294373	1531019	2902397	-1371378	631687	554682	77005
2011	2303466	3603059	-1299593	1737678	3104453	-1366775	565788	498606	67182
2012	2450164	3537127	-1086983	1880663	3029539	-1148876	569501	507588	61913

year

delta

There are a total of 10 fields of information in each line. Our programming objective uses only the first and fourth fields, which are arbitrarily called "year" and "delta" respectively. We will ignore all the other fields of data.

Designing and Implementing the Mapper Class



Input to the Mapper Class

```
input format = TextInputFormat  
Key = LongWritable  
Value = Text
```

} More detail later

First record:

input key input value

8 1901 988 925 63 988 925 63

```
StringTokenizer itr = new StringTokenizer(value.toString(),"\\s+");
```

Let us define and implement the Mapper class to solve the programming problem.

In this example, we will be using the TextInputFormat class as the type of input to the Mapper class. Using this format, the key from the default record reader associated with TextInputFormat is the byte offset into the file (LongWritable). We won't be using this key for anything in our program, so we will just ignore it. The value from the default record reader is the line read from the input file, in Text.



Output of the Mapper Class

output key

output value

```
context.write(new Text("summary"), new Text(year + "_" + delta));
```

First few lines of output

summary	1901_63
summary	1902_77
summary	1903_45
summary	1904_-43
summary	1905_-23
summary	1906_25

After grabbing fields 1 and 4 from a record, the Mapper class emits the constant key ("summary") and a composite output value of the year (field 1) followed by an underscore followed by the delta (field 4).



Design and Implement the Mapper Class

```
public class ReceiptsMapper extends Mapper  
<LongWritable,Text,Text> {  
    public void map(LongWritable key, Text value, Context context)  
throws IOException, InterruptedException {  
    StringTokenizer iterator = new  
StringTokenizer(value.toString(),"\\s+");  
    String year = iterator.nextToken();  
    iterator.nextToken();  
    iterator.nextToken();  
    String delta= iterator.nextToken();  
    context.write(new Text("summary"), new Text(year + " - "  
delta));  
}  
}
```

Phases of Developing a MapReduce Application

- Configuration API
- Configuring the Development Environment
- GenericOptionsParser, Tool and ToolRunner
- Writing Unit Tests
- Running locally and in a cluster on Test Data
- The MapReduce Web UI
- Hadoop Logs
- Tuning a Job to improve performance

Stages 1:Developing a MapReduce Application

- Writing a program in MapReduce follows a certain pattern.
- You start by writing your map and reduce functions, ideally with unit tests to make sure they do what you expect.
- Then you write a driver program to run a job, which can run from your IDE using a small subset of the data to check that it is working.
- If it fails, you can use your IDE's debugger to find the source of the problem.
- With this information, you can expand your unit tests to cover this case and improve your mapper or reducer as appropriate to handle such input correctly.

Stages 2:Developing a MapReduce Application

- When the program runs as expected against the small dataset, you are ready to unleash it on a cluster.
- Running against the full dataset is likely to expose some more issues, which you can fix as before, by expanding your tests and mapper or reducer to handle the new cases.
- Debugging failing programs in the cluster is a challenge, so we look at some common techniques to make it easier.

Stage 3: Developing a MapReduce Application

- After the program is working, you may wish to do some tuning, first by running through some standard checks for making MapReduce programs faster and then by doing task profiling.
- Profiling distributed programs is not easy, but Hadoop has hooks to aid the process.

Example: Word Count

Task: Counting the word occurrences (frequencies) in a text file (or set of files).

`<word, count >as <key, value >pair`

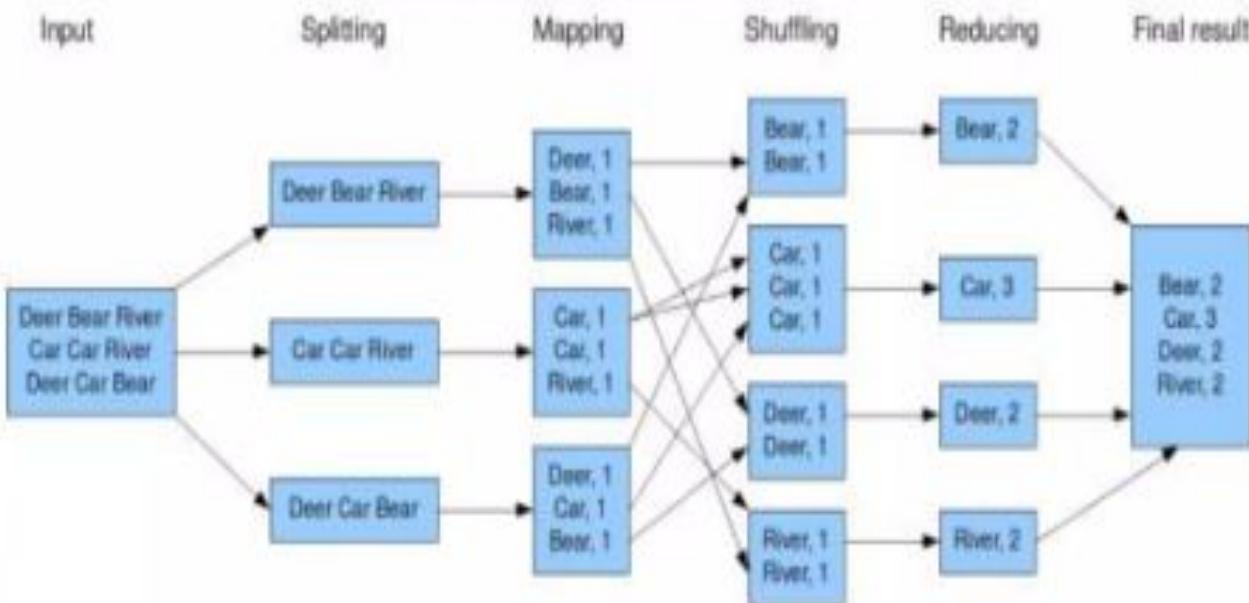
Mapper: Emits `<word, 1 >`for each word (no counting at this part).

Shuffle in between: pairs with same keys grouped together and passed to a single machine.

Reducer: Sums up the values (1s) with the same key value

Example: Word Count

The overall MapReduce word count process



Example: Job Tracker

Tasks

Hadoop map task list for job_200904110811_0003 on ip-10-250-110-47

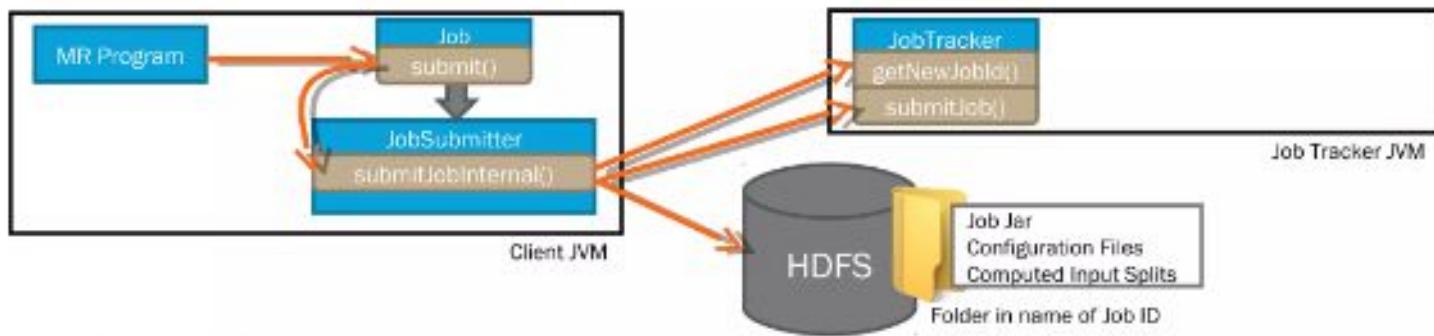
Completed Tasks

Task	Complete	Status	Start Time	Finish Time	Errors	Counters
task_200904110811_0003_m_000043	100.00%	hdfs://ip-10-250-110-47.ec2.internal/user/root/input/hcdcl/all/1949.gz:0+220338475	11-Apr-2009 09:00:06	11-Apr-2009 09:01:25 (1mins, 18sec)		10
task_200904110811_0003_m_000044	100.00%	Detected possibly corrupt record: see logs.	11-Apr-2009 09:00:06	11-Apr-2009 09:01:28 (1mins, 21sec)		11
task_200904110811_0003_m_000045	100.00%	hdfs://ip-10-250-110-47.ec2.internal/user/root/input/hcdcl/all/1970.gz:0+208374610	11-Apr-2009 09:00:06	11-Apr-2009 09:01:28 (1mins, 21sec)		10



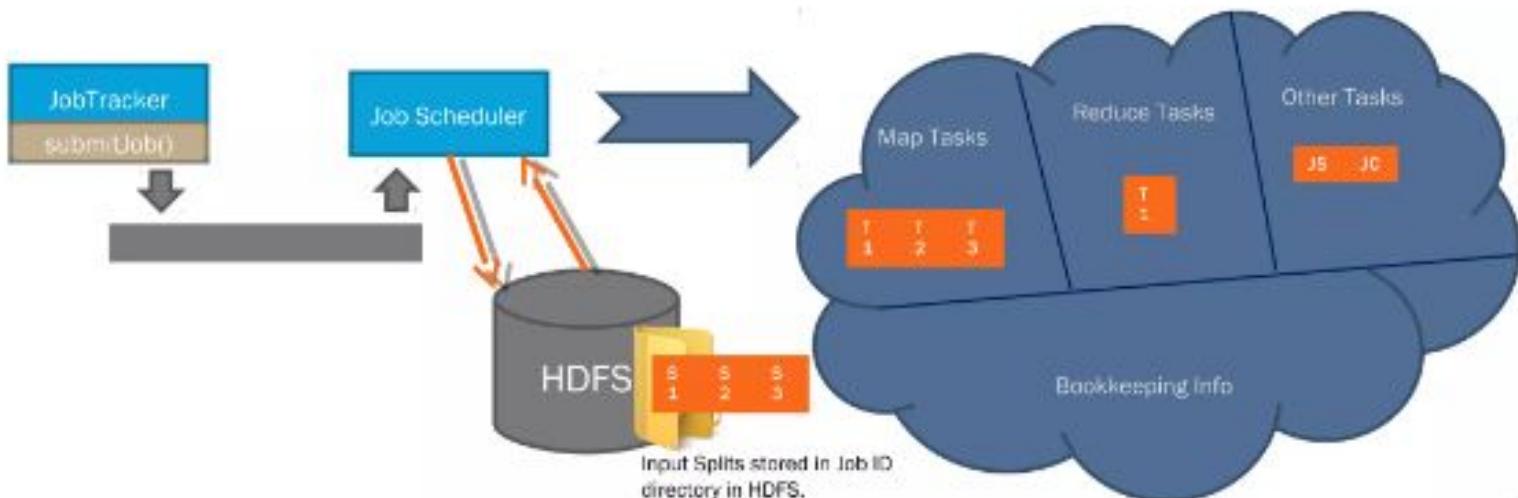
ANATOMY OF CLASSIC MAPREDUCE IN HADOOP

FIRST A JOB HAS TO BE SUBMITTED TO HADOOP CLUSTER. LET'S SEE HOW JOB SUBMISSION HAPPENS IN HADOOP.



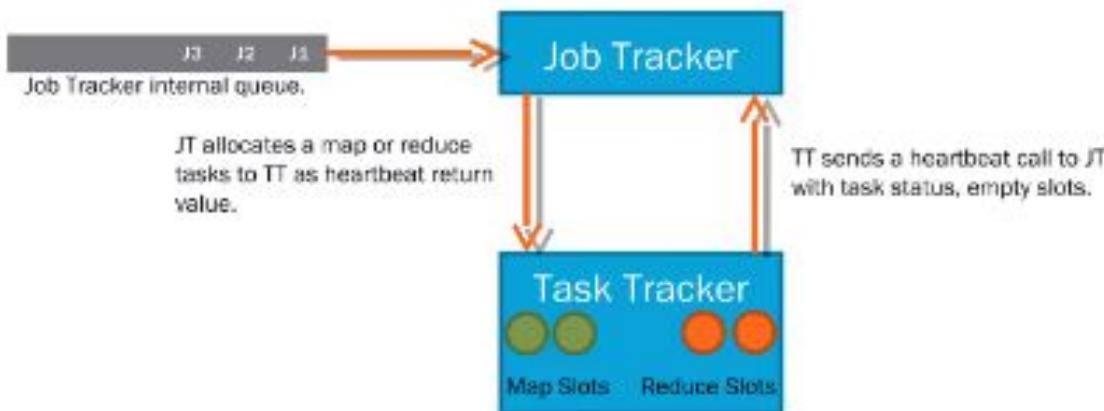
- Client calls the submit() (or waitForCompletion()) method on Job.
- Job creates a new instance of JobSubmitter which has a method called 'submitJobInternal()'.
- Then the Job.submit() method calls JobSubmitter.submitJobInternal() method which does the following;
 - Invokes the JobTracker.getNewJobId() to generate a unique id for the job.
 - Checks the output specification. If output path is not specified or if it already exists, then an exception is thrown.
 - Checks the input specification. In case of invalid input path or if splits couldn't be computed, then an exception is thrown. Else, the input splits are computed.
 - Creates a directory in the name of Job Id in HDFS.
 - Copies Job jar, configuration files and computed input splits to this directory. The Job jar has **high replication factor which is configured in 'mapred.submit.replication' property.**
 - Informs the JobTracker that the job is ready to be executed by calling submitJob() on JobTracker.
- Instead of Job.submit() the MapReduce program can call Job.waitForCompletion() method. Difference here is, the later waits till the job is complete by polling the JobTracker for every second. Once the job is completed successfully, counters are displayed in console. In case of job failure, the exception is displayed in console. But the submit() method terminates once the job submission is done.

NEXT THE SUBMITTED JOB WILL BE INITIALIZED.
NOW LET'S SEE HOW JOB INITIALIZATION
HAPPENS IN HADOOP.



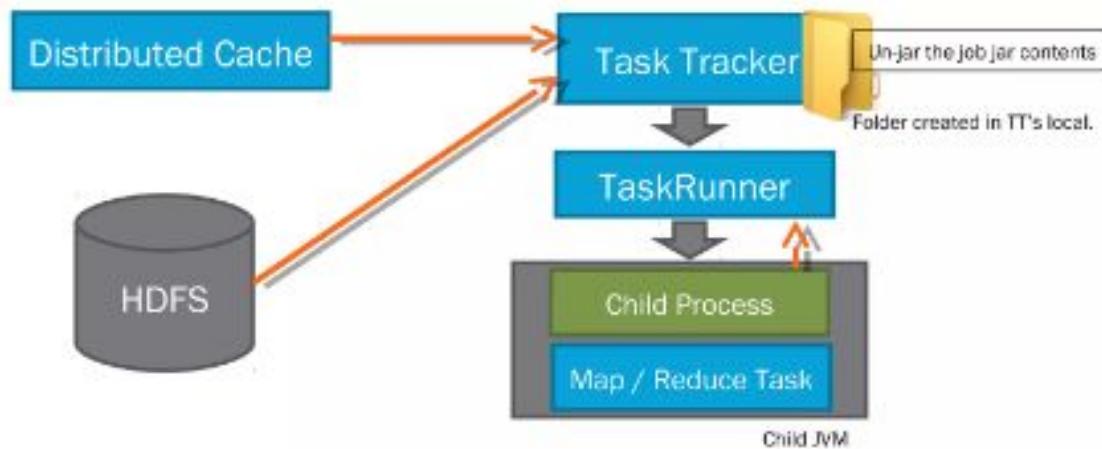
- When `JobTracker.submitJob()` is called, `JobTracker` adds a token to an internal queue.
- A job scheduler picks it up and creates an object representation for the job. This representation encapsulates tasks & bookkeeping infos (to track the status & progress of job's tasks).
- Job scheduler then reads the input splits from HDFS.
- Then the scheduler creates one map task for each input split.
- '`mapred.reduce.tasks`' property is meant to have an integer. Job scheduler reads this property and creates those many number of reduce tasks. Let's assume the property's value was 1.
- Job setup (JS) & Job cleanup (JC) are the other 2 tasks created. Job setup task is run before any map task and Job cleanup task is run after all reduce tasks are complete.
- Each job has an `OutputCommitter` (a Java Interface) associated with it. Default implementation is `FileOutputCommitter`. The `OutputCommitter` defines what setup and cleanup task (for job & task) should do.
- In case of `FileOutputCommitter`, the job setup task will create the final output directory for the job and temp working space for task output. For cleanup task, it deletes the temp working space for task output. Please refer the API document for more info.

INITIALIZED JOB IS SPLIT INTO TASKS AND JOB
TRACKER ASSIGNS TASKS TO TASK TRACKER.
NOW LET'S SEE HOW TASK ASSIGNMENT
HAPPENS IN HADOOP.



- Based on factors like available memory, available cores, etc., fixed number of map and reduce slots are configured in a Task Tracker.
- When Task Tracker is up and running, it sends a heartbeat call (for every 5 seconds) to Job Tracker. This is a two way communication between Task Tracker & Job Tracker. Task Tracker uses this call to inform Job Tracker that its alive. Also Task Tracker uses this call to inform Job Tracker about the status of tasks and which all map & reduce slots are available.
- Job Tracker might have many jobs in queue to run. It uses one of the scheduling algorithms to pick a job from queue. Once a job is picked, its associated tasks will be picked.
- When Job Tracker finds that there are empty slots (from the heartbeat call), it allocates tasks to Task Tracker using the heartbeat return value.
- Job Tracker considers the data locality when allocating map tasks; i.e., Job Tracker tries to allocate the map job to a Task Tracker where the block is available (which is called data local). If its not possible, Job Tracker tries to allocate the map task to a map slot in same rack (which is called rack local). There are no such considerations for reduce task.

NOW TASKS ARE ASSIGNED TO TASK TRACKER WHICH FOLLOWS A SERIES OF STEPS TO EXECUTE A TASK. LET'S SEE HOW TASKS ARE EXECUTED IN TASK TRACKER.



- Now the Task Tracker has been assigned a task.
- Task Tracker copies the job jar from HDFS to task tracker's file system. Also it copies required files from Distributed Cache.
- Task Tracker creates a new folder in task tracker's file system.
- Job jar content is un-jared into the new folder.
- Creates a new instance of TaskRunner.
- TaskRunner launches a new JVM to run each task, so that any bug in the user-defined map & reduce functions don't affect the Task Tracker.**
- The child process communicates with its parent through the umbilical interface. It informs the parent of the task's progress every few seconds until the task is complete.
- There is a setup and cleanup tasks executed in the same JVM where the task is executed. The OutputCommitter implementation associated with the job determines what action to be taken during startup and cleanup.

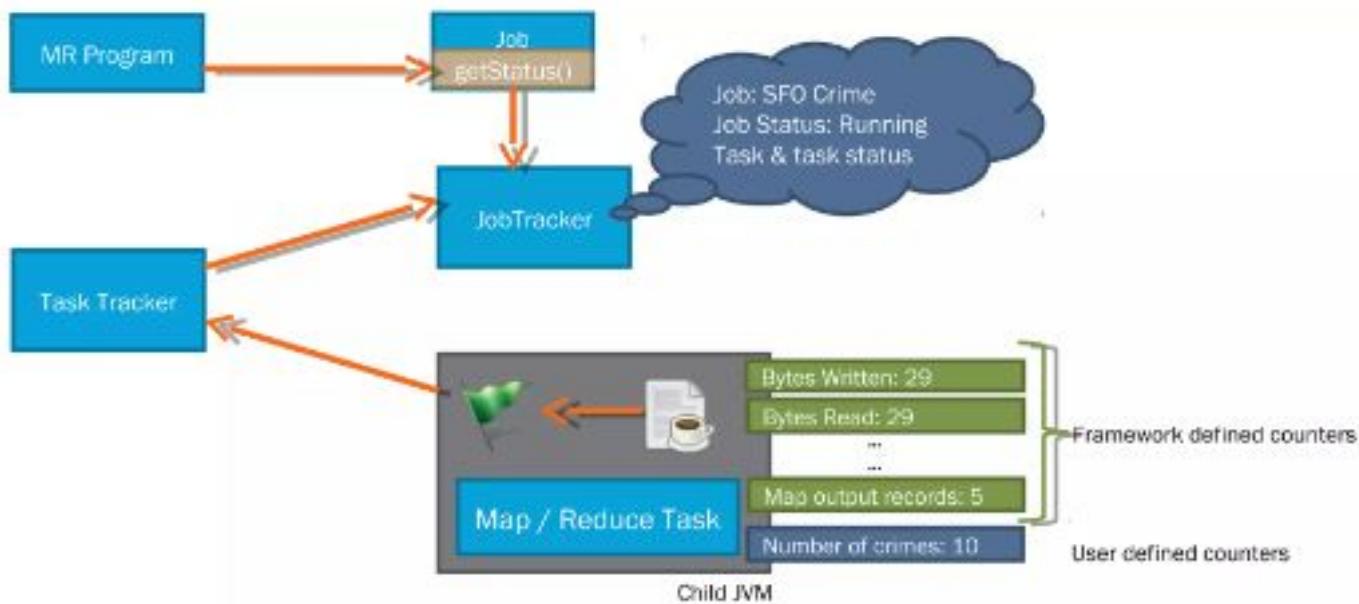
SINCE TASKS ARE EXECUTED IN A DISTRIBUTED ENVIRONMENT, TRACKING THE PROGRESS AND STATUS OF JOB IS TRICKY. LET'S SEE HOW PROGRESS AND STATUS UPDATES ARE TAKEN CARE IN HADOOP.

- MapReduce jobs are long running batch job which takes anything from one minute to hours to run. Because of this user has to get feedback on how job is processing.
- Each job and task have status which comprises of the following:

Status of Job or Task	Progress of map & reduce	Value of job counters	Status description
-----------------------	--------------------------	-----------------------	--------------------

Job/Task Status

- Status of Job/Task: Possible values are RUNNING, SUCCESSFULLY COMPLETED and FAILED. Job tracker and Task trackers set the value as the job or task progresses.
- Each task tracks the program by tracking the proportion of the task completed.
 - Map task's progress is measured by the proportion of input processed so far.
 - Measuring Reduce task's progress is little tricky. The system does it by dividing the total progress into 3 parts corresponding to 3 phases of Shuffle.

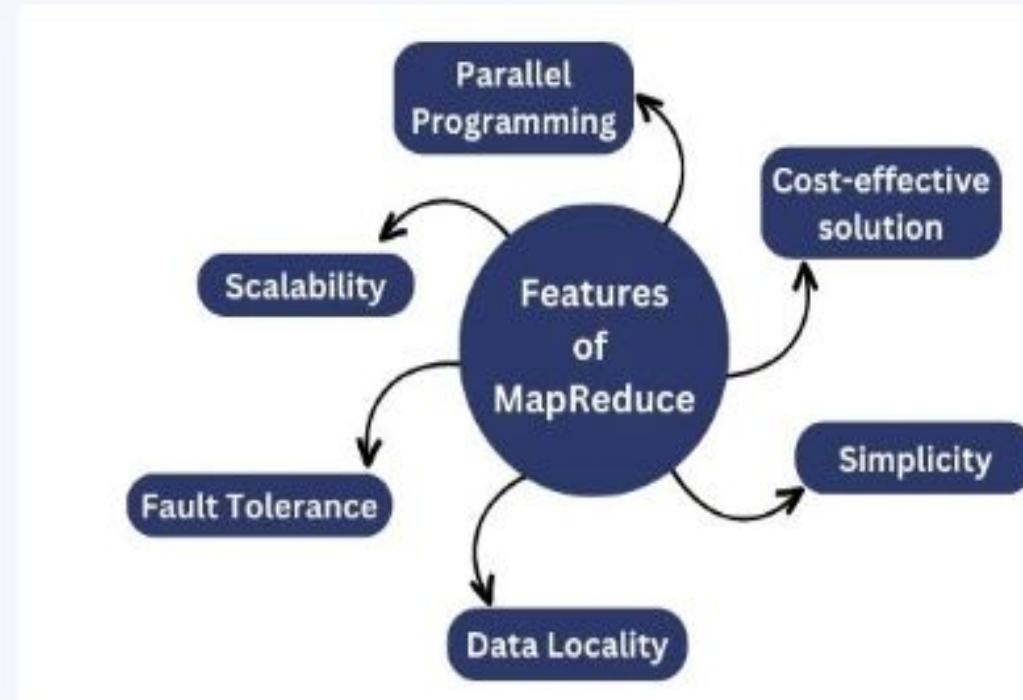


- Each task has set of counters that count various events as the task runs.
- Most of these counters are build into the framework. We can define our own counters. These are called user defined counters.
- As the tasks progress, it sets a flag to indicate that its time to send the progress to Task Tracker.
- This flag is checked every second by a thread and notifies the Task Tracker.
- For every 5 seconds, the Task Tracker sends a heart beat to Job Tracker. The status of all tasks are sent to Job tracker along with the heartbeat call.
- Counters are sent less frequently than every 5 seconds because they can be relatively high-bandwidth.
- Job Tracker combines these updates to produce a global view of the status of all the jobs & its related tasks.
- Clients invoke the API `Job.getStatus()` for every second to get the status from Job Tracker.

THIS EXECUTION PROCESS CONTINUES TILL ALL THE TASKS ARE COMPLETED. ONCE THE LAST TASK IS COMPLETED, MR FRAMEWORK ENTERS THE LAST PHASE CALLED JOB COMPLETION.

Key Features of MapReduce

There are some key features of MapReduce below:



Scalability

MapReduce can scale to process vast amounts of data by distributing tasks across a large number of nodes in a cluster. This allows it to handle massive datasets, making it suitable for Big Data applications.

Fault Tolerance

MapReduce incorporates built-in fault tolerance to ensure the reliable processing of data. It automatically detects and handles node failures, rerunning tasks on available nodes as needed.

Data Locality

MapReduce takes advantage of data locality by processing data on the same node where it is stored, minimizing data movement across the network and improving overall performance.

Simplicity

The MapReduce programming model abstracts away many complexities associated with distributed computing, allowing developers to focus on their data processing logic rather than low-level details.

Cost-Effective Solution

Hadoop's scalable architecture and MapReduce programming framework make storing and processing extensive data sets very economical.

Parallel Programming

Tasks are divided into programming models to allow for the simultaneous execution of independent operations. As a result, programs run faster due to parallel processing, making it easier for a process to handle each job. Thanks to parallel processing, these distributed tasks can be performed by multiple processors. Therefore, all software runs faster.

Features explained in detail:

1. Highly scalable

A framework with excellent scalability is Apache Hadoop MapReduce. This is because of its capacity for distributing and storing large amounts of data across numerous servers. These servers can all run simultaneously and are all reasonably priced.

By adding servers to the cluster, we can simply grow the amount of storage and computing power. We may improve the capacity of nodes or add any number of nodes (horizontal scalability) to attain high computing power. Organizations may execute applications from massive sets of nodes, potentially using thousands of terabytes of data, thanks to Hadoop MapReduce programming.

2. Versatile

Businesses can use MapReduce programming to access new data sources. It makes it possible for companies to work with many forms of data. Enterprises can access both organized and unstructured data with this method and acquire valuable insights from the various data sources.

Since Hadoop is an open-source project, its source code is freely accessible for review, alterations, and analyses. This enables businesses to alter the code to meet their specific needs. The MapReduce framework supports data from sources including email, social media, and clickstreams in different languages.

3. Secure

The MapReduce programming model uses the HBase and HDFS security approaches, and only authenticated users are permitted to view and manipulate the data. HDFS uses a replication technique in Hadoop 2 to provide fault tolerance. Depending on the replication factor, it makes a clone of each block on the various machines. One can therefore access data from the other devices that house a replica of the same data if any machine in a cluster goes down. Erasure coding has taken the role of this replication technique in Hadoop 3. Erasure coding delivers the same level of fault tolerance with less area. The storage overhead with erasure coding is less than 50%.

4. Affordability

With the help of the MapReduce programming framework and Hadoop's scalable design, [big data](#) volumes may be stored and processed very affordably. Such a system is particularly cost-effective and highly scalable, making it ideal for business models that must store data that is constantly expanding to meet the demands of the present.

In terms of scalability, processing data with older, conventional relational database management systems was not as simple as it is with the Hadoop system. In these situations, the company had to minimize the data and execute classification based on presumptions about how specific data could be relevant to the organization, hence deleting the raw data. The MapReduce programming model in the Hadoop scale-out architecture helps in this situation.

5. Fast-paced

The Hadoop Distributed File System, a distributed storage technique used by MapReduce, is a mapping system for finding data in a cluster. The data processing technologies, such as MapReduce programming, are typically placed on the same servers that enable quicker data processing.

Thanks to Hadoop's distributed data storage, users may process data in a distributed manner across a cluster of nodes. As a result, it gives the Hadoop architecture the capacity to process data exceptionally quickly. Hadoop MapReduce can process unstructured or semi-structured data in high numbers in a shorter time.

6. Based on a simple programming model

Hadoop MapReduce is built on a straightforward programming model and is one of the technology's many noteworthy features. This enables programmers to create MapReduce applications that can handle tasks quickly and effectively. Java is a very well-liked and simple-to-learn programming language used to develop the MapReduce programming model.

Java programming is simple to learn, and anyone can create a data processing model that works for their company. Hadoop is straightforward to utilize because customers don't need to worry about computing distribution. The framework itself does the processing.

7. Parallel processing-compatible

The parallel processing involved in MapReduce programming is one of its key components. The tasks are divided in the programming paradigm to enable the simultaneous execution of independent activities. As a result, the program runs faster because of the parallel processing, which makes it simpler for the processes to handle each job. Multiple processors can carry out these broken-down tasks thanks to parallel processing. Consequently, the entire software runs faster.

8. Reliable

The same set of data is transferred to some other nodes in a cluster each time a collection of information is sent to a single node. Therefore, even if one node fails, backup copies are always available on other nodes that may still be retrieved whenever necessary. This ensures high data availability.

The framework offers a way to guarantee data trustworthiness through the use of Block Scanner, Volume Scanner, Disk Checker, and Directory Scanner modules. Your data is safely saved in the cluster and is accessible from another machine that has a copy of the data if your device fails or the data becomes corrupt.

9. Highly available

Hadoop's fault tolerance feature ensures that even if one of the DataNodes fails, the user may still access the data from other DataNodes that have copies of it. Moreover, the high accessibility Hadoop cluster comprises two or more active and passive NameNodes running on hot standby. The active NameNode is the active node. A passive node is a backup node that applies changes made in active NameNode's edit logs to its namespace.